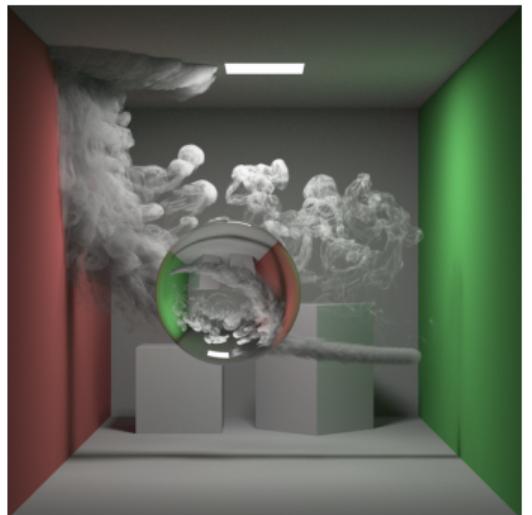
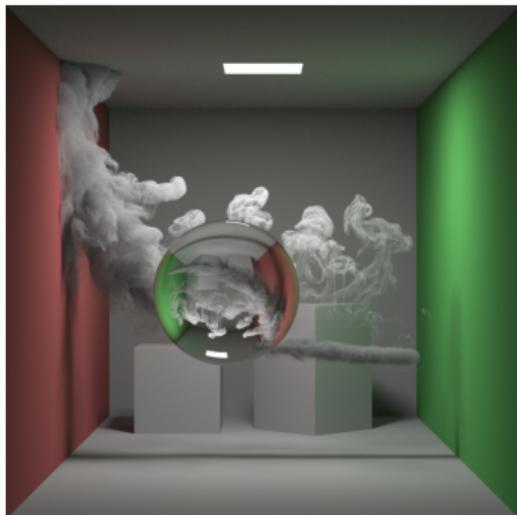
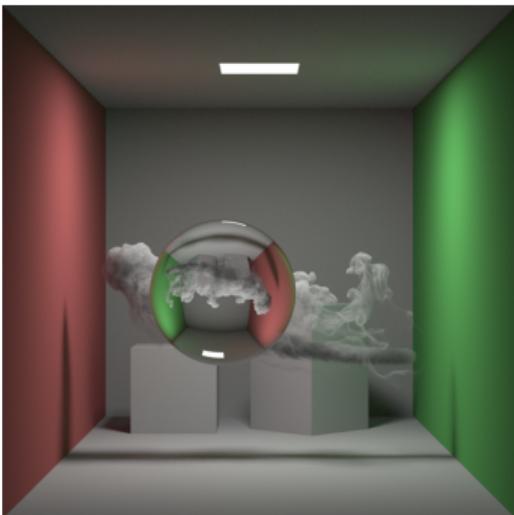


# Decoupled Ray-marching of Heterogeneous Participating Media

Christopher Kulla



## Goals

- ▶ Add volume rendering to our production raytracer (smoke, clouds, fire,...)
- ▶ Support arbitrary shaders (including time-varying)
- ▶ Avoid light caches, preprocessing, and large data structures
- ▶ Take advantage of previous work on lighting homogeneous media

## Challenge

- ▶ Heterogeneous media add a *nested* integral to compute transmission
- ▶ This integral is *dependent* on the sample position along the line

### Volume Rendering Equation [Kajiya and Von Herzen, 1984]

$$L(x, \vec{\omega}) = \int_a^b \sigma_s(x + t\vec{\omega}) e^{-\tau(t)} \left( \int_{4\pi} \rho(\vec{\omega}, \vec{\omega}') L(x + t\vec{\omega}, \vec{\omega}') d\vec{\omega}' \right) dt$$
$$\tau(t) = \int_0^t \sigma_t(x + t'\vec{\omega}) dt'$$

## Previous Work: Unbiased Methods

- ▶ “Unbiased Global Illumination with Participating Media” [Raab et al, 2006]
- ▶ “Image Synthesis using Adjoint Photons” [Morley et al, 2006]
- ▶ “Unbiased, Adaptive Stochastic Sampling for Rendering Inhomogeneous Participating Media” [Yue et al, 2010]

## Unbiased Methods

- ▶ Use rejection sampling to implicitly generate samples according to  $e^{-\tau(t)}$
- ▶ Requires an upper bound on  $\sigma_t$  for the ray (procedural shaders?)
- ▶ Acceleration structure [Yue et al] increases memory usage (motion blur?)
- ▶ Sample distribution is implicit, pdf is not directly evaluable

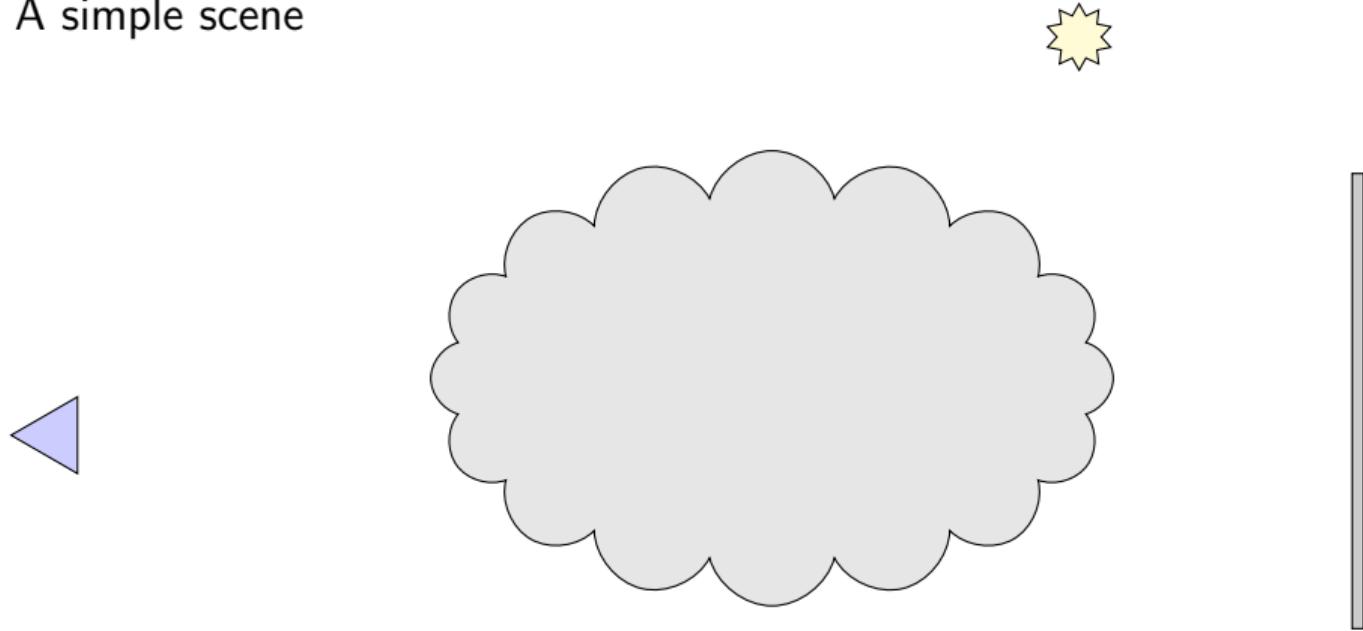
## Previous Work: Biased Methods (1/2)

Computing lighting at every sample:

- ▶ “Hypertexture” [Perlin and Hoffert, 1989]
- ▶ “Production Volume Rendering” [Course notes, 2011]
- ▶ Ubiquitous method for VFX volume rendering

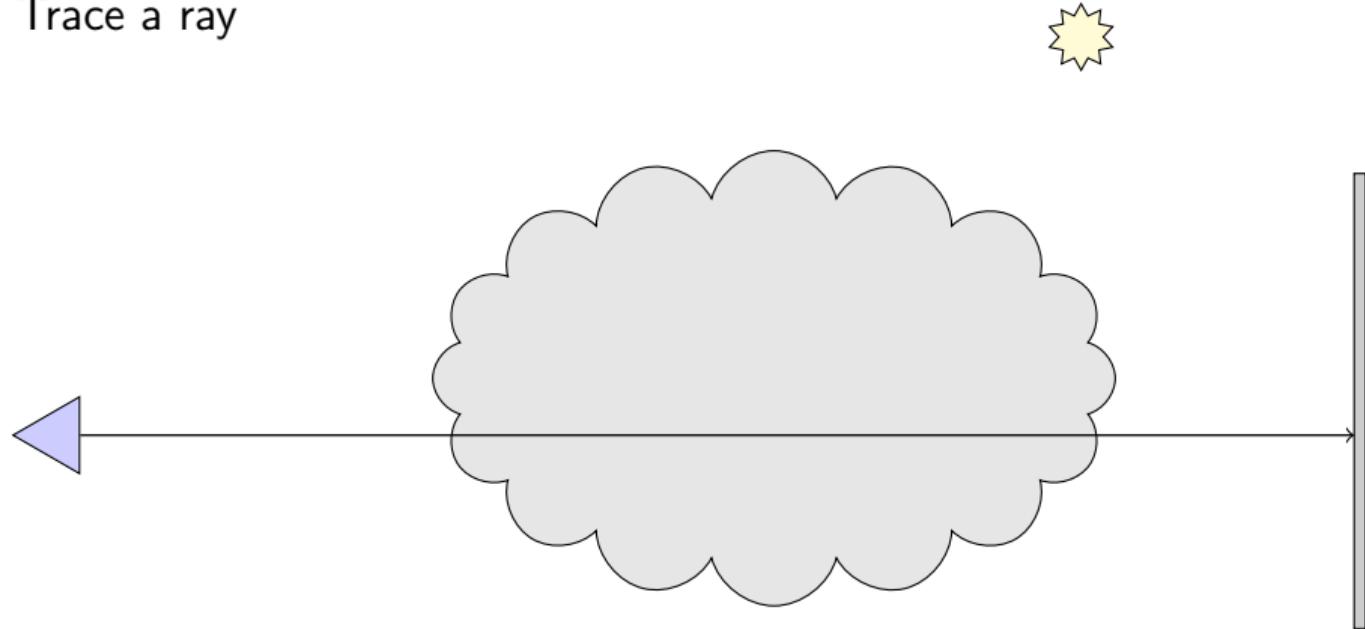
# Traditional Ray-marching

A simple scene



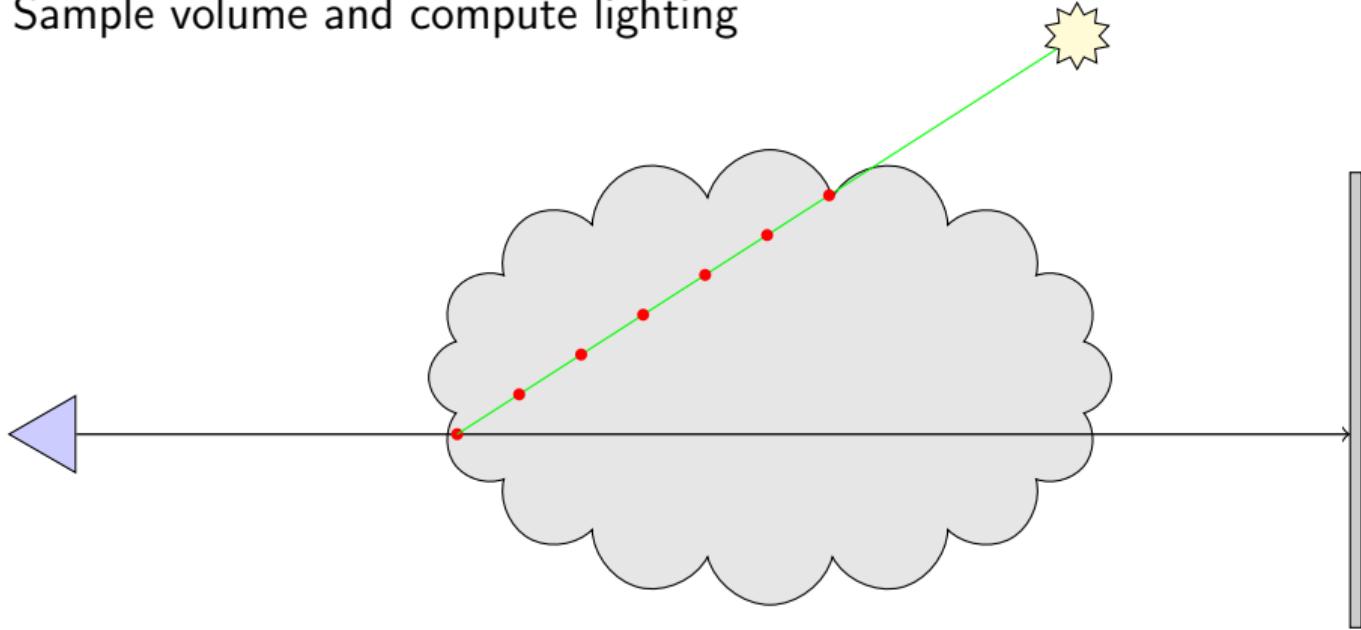
# Traditional Ray-marching

Trace a ray



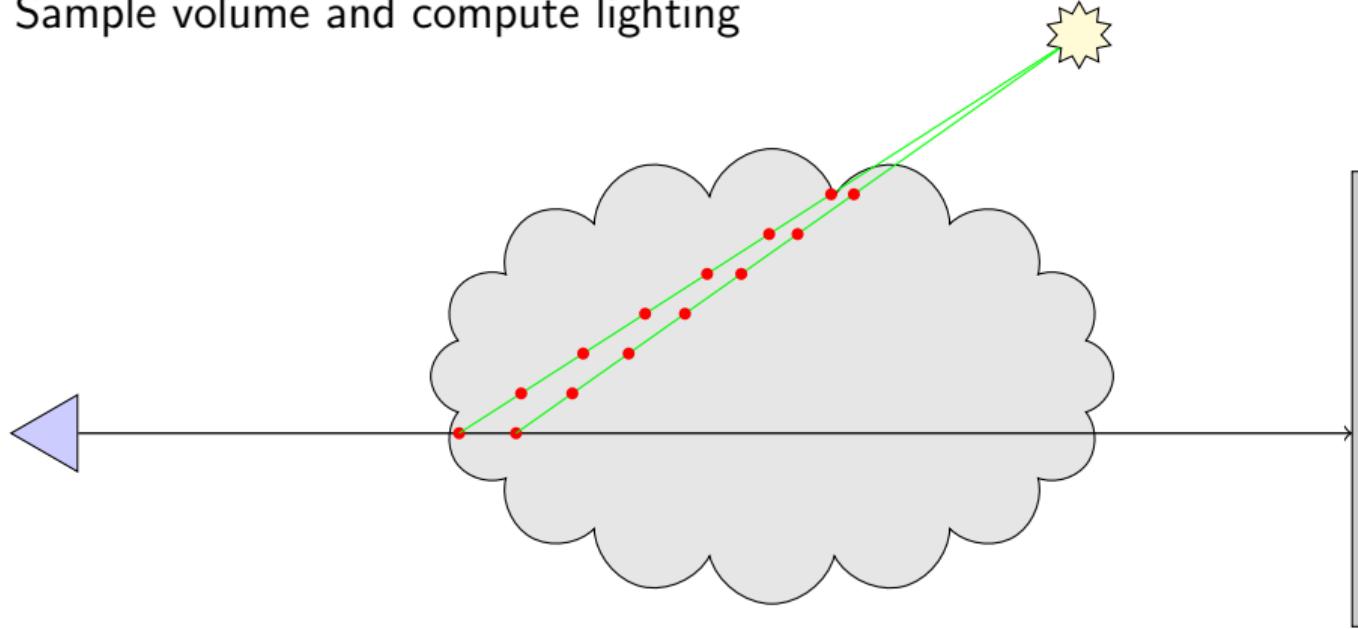
# Traditional Ray-marching

Sample volume and compute lighting



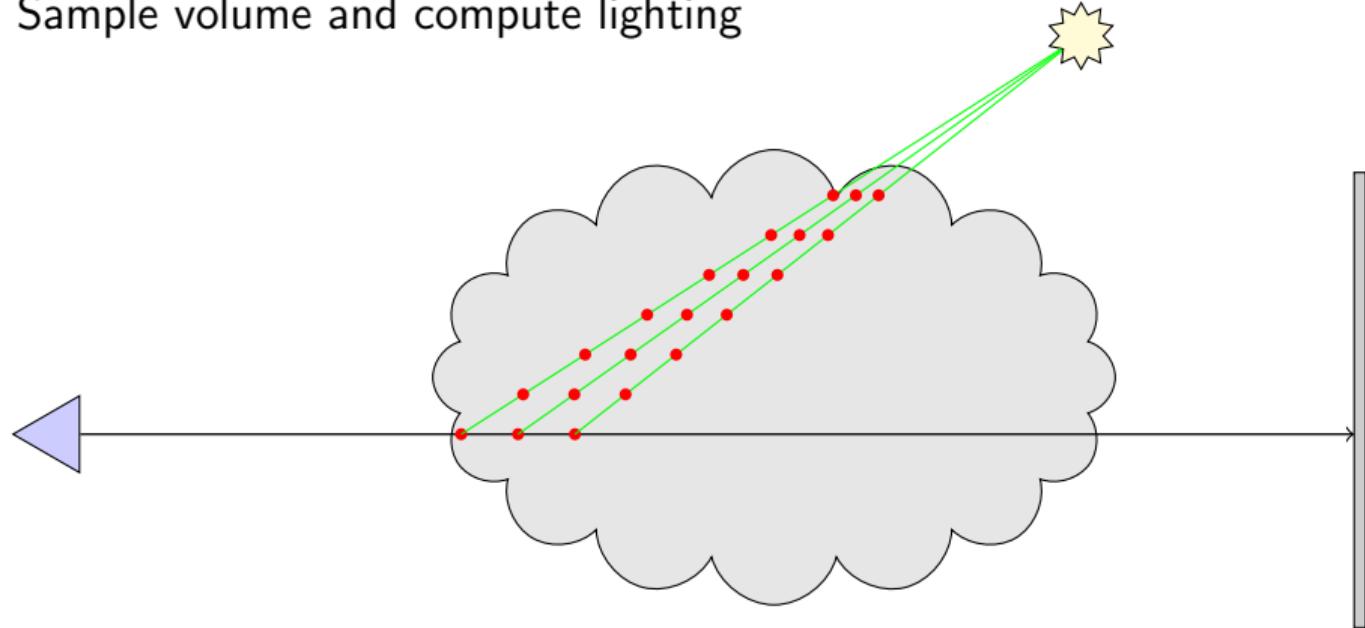
# Traditional Ray-marching

Sample volume and compute lighting



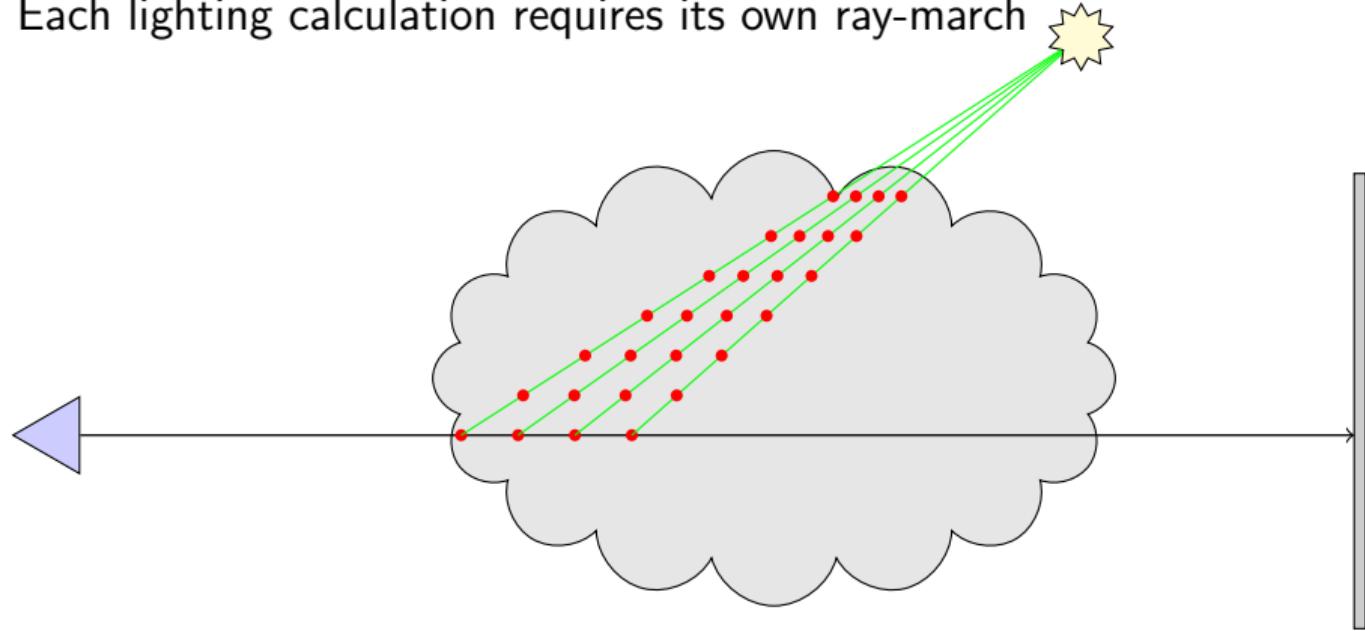
# Traditional Ray-marching

Sample volume and compute lighting



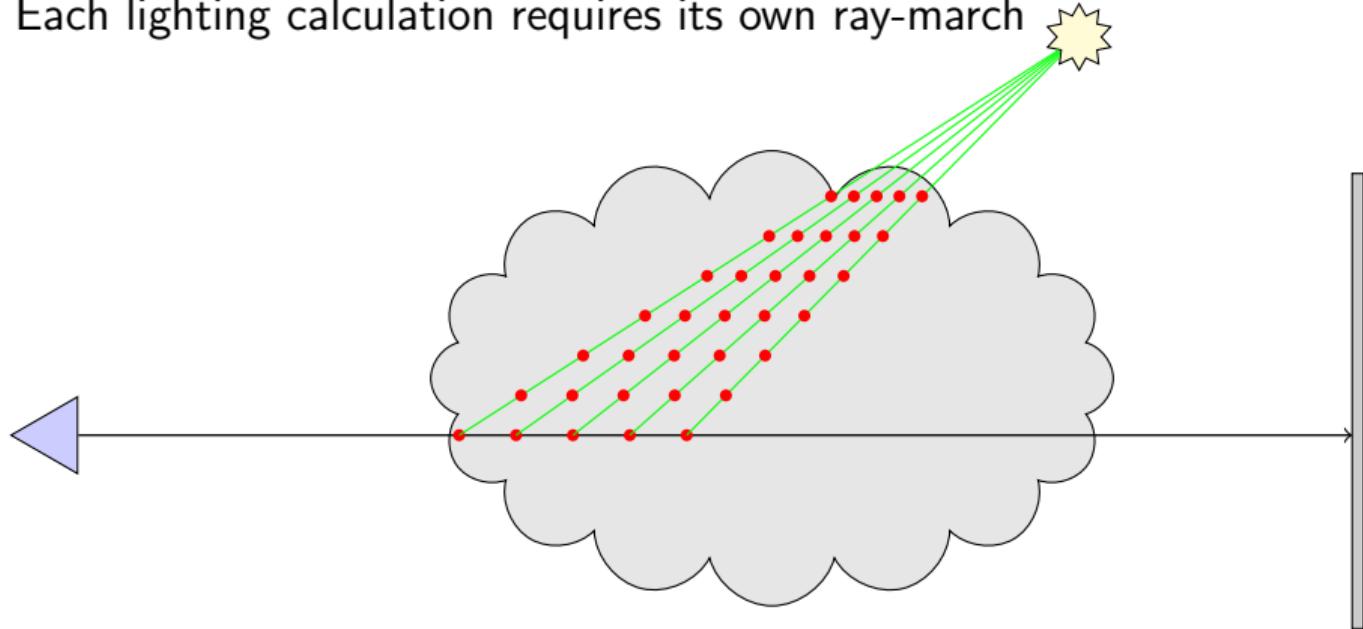
# Traditional Ray-marching

Each lighting calculation requires its own ray-march



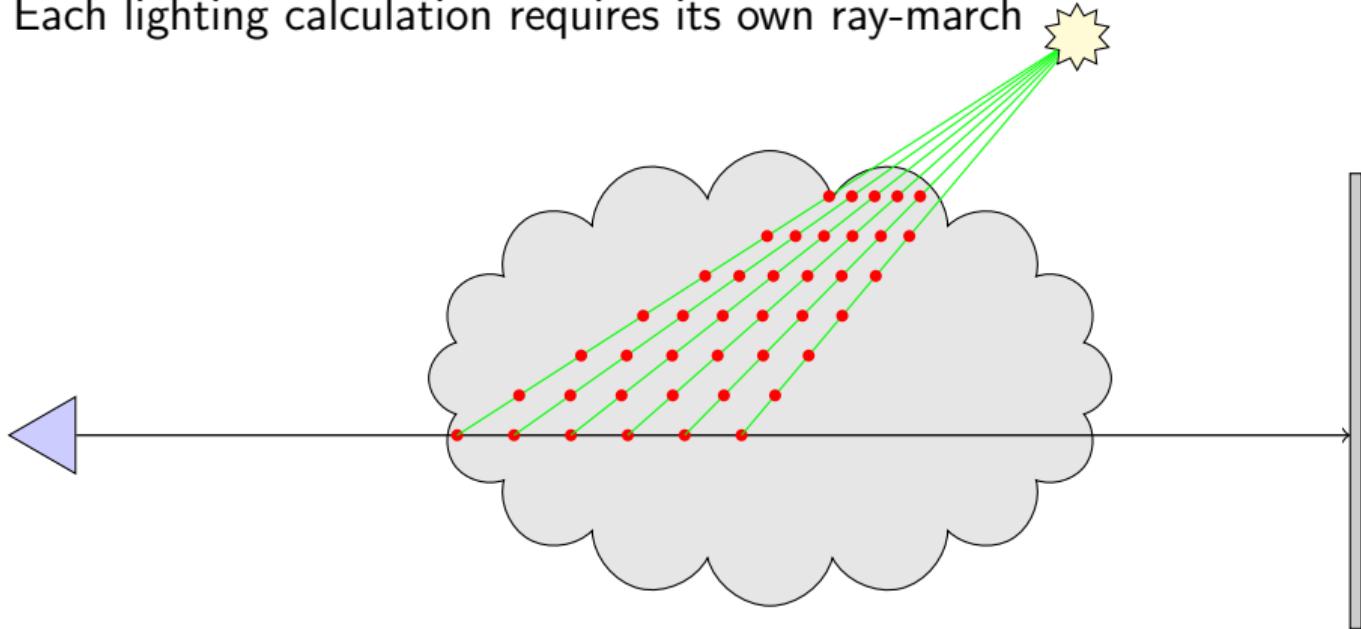
# Traditional Ray-marching

Each lighting calculation requires its own ray-march



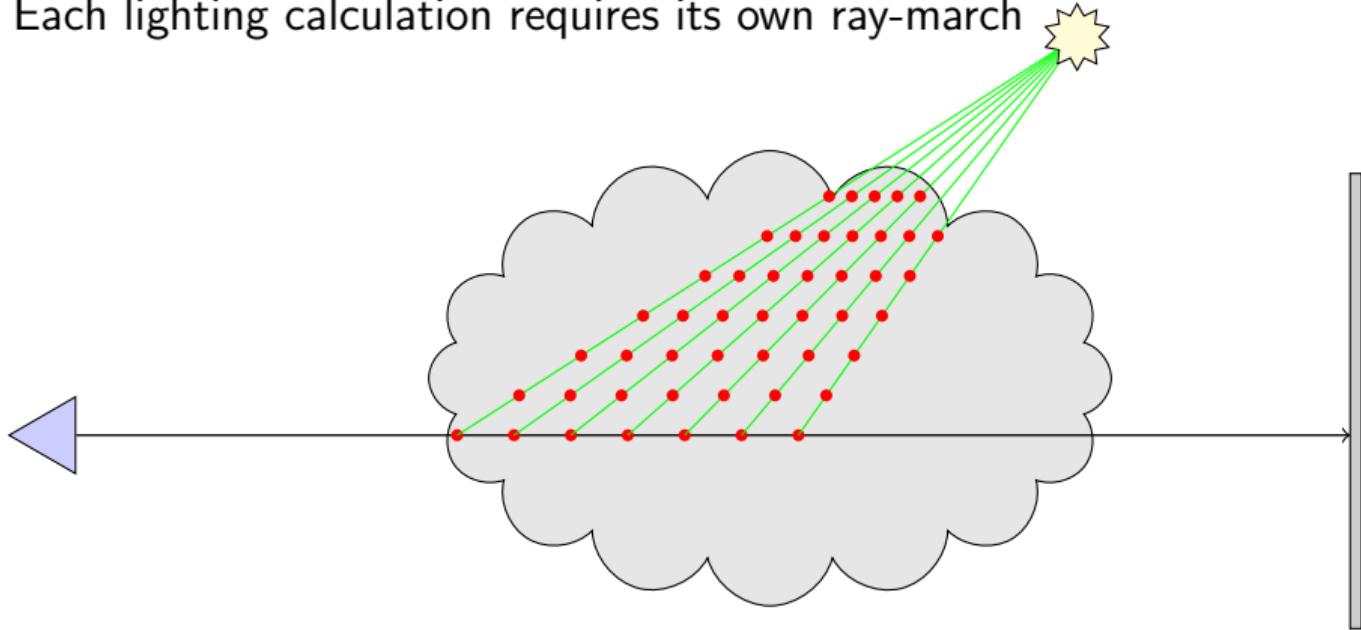
# Traditional Ray-marching

Each lighting calculation requires its own ray-march



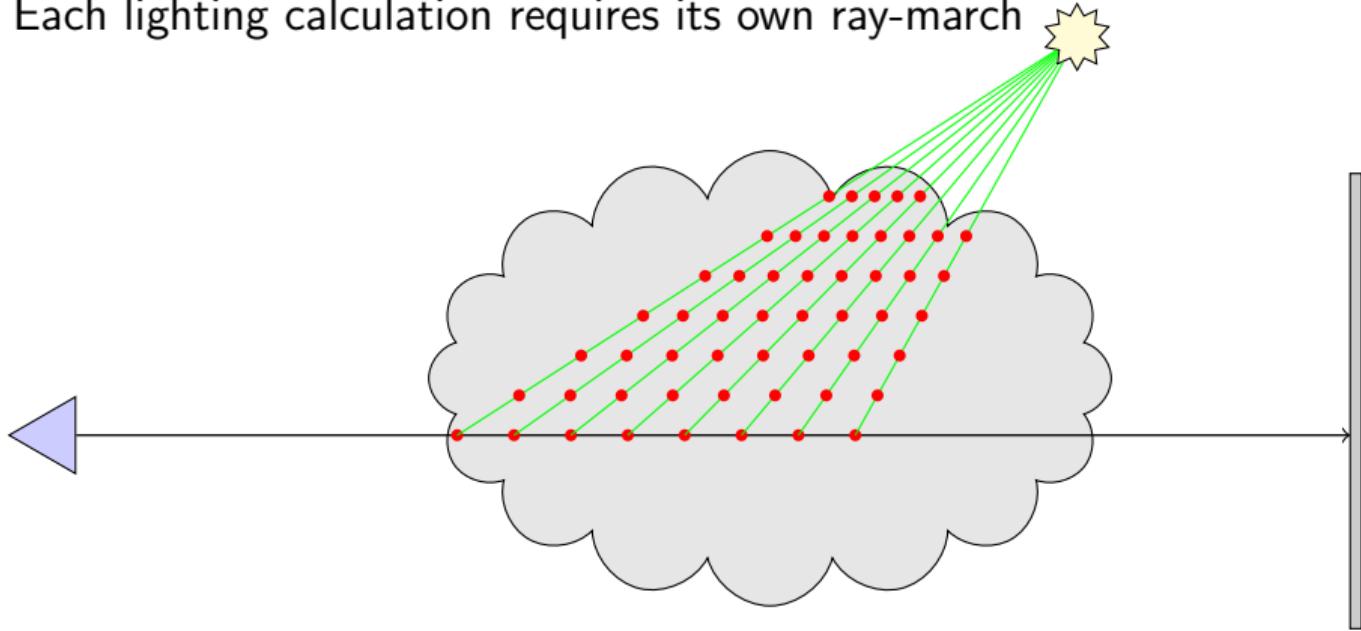
# Traditional Ray-marching

Each lighting calculation requires its own ray-march



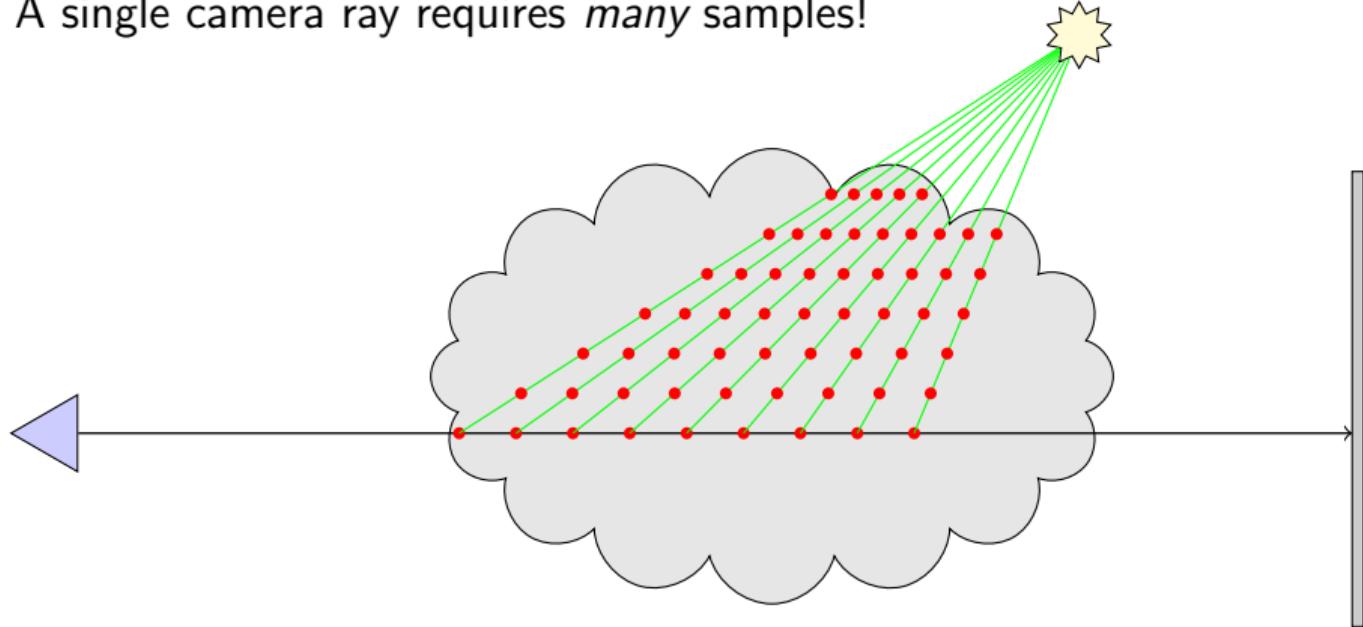
# Traditional Ray-marching

Each lighting calculation requires its own ray-march



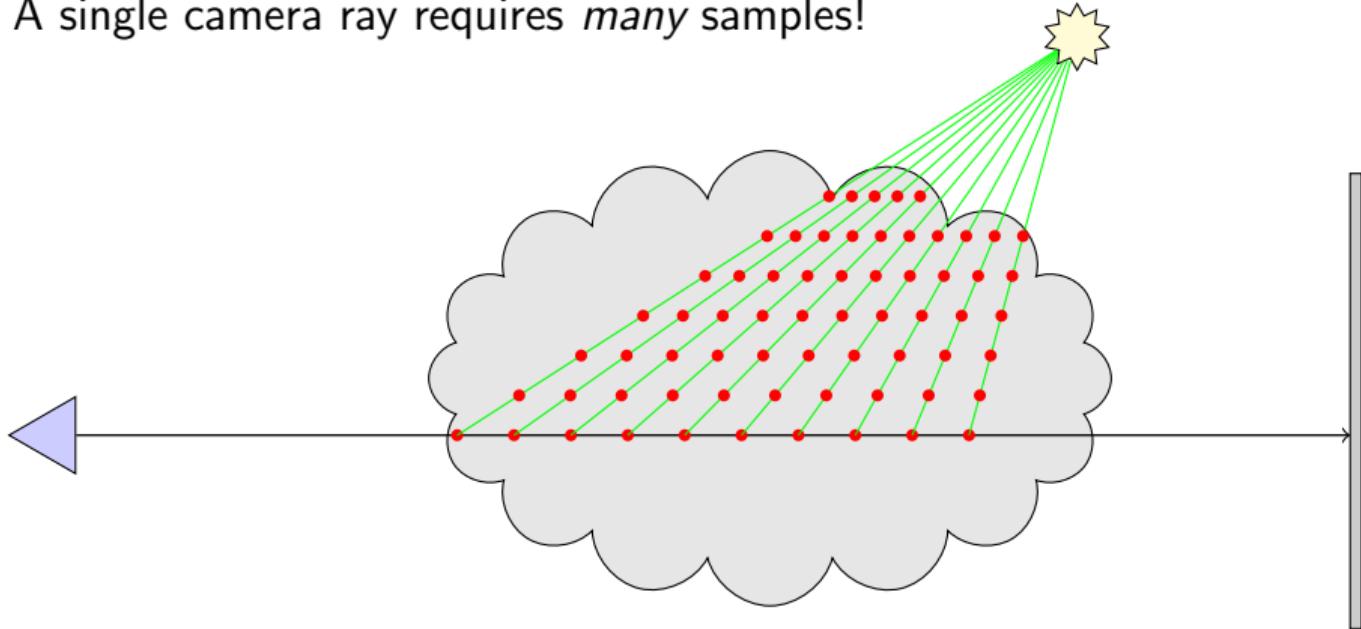
# Traditional Ray-marching

A single camera ray requires *many* samples!



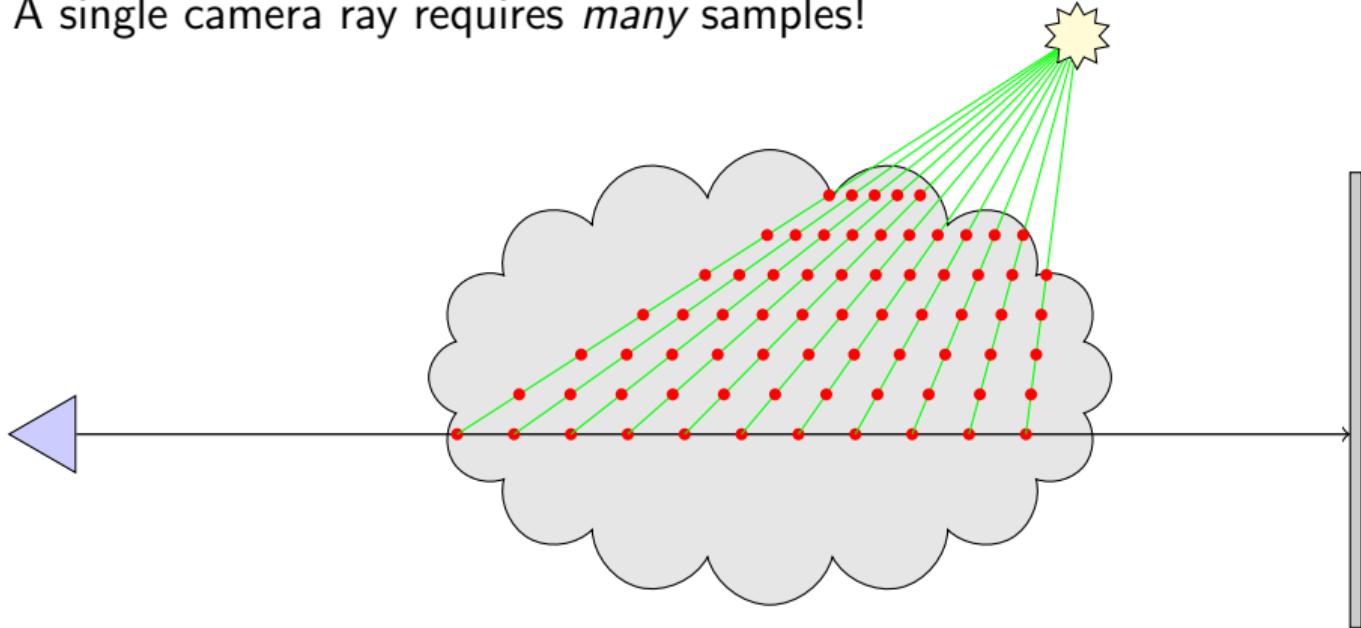
# Traditional Ray-marching

A single camera ray requires *many* samples!



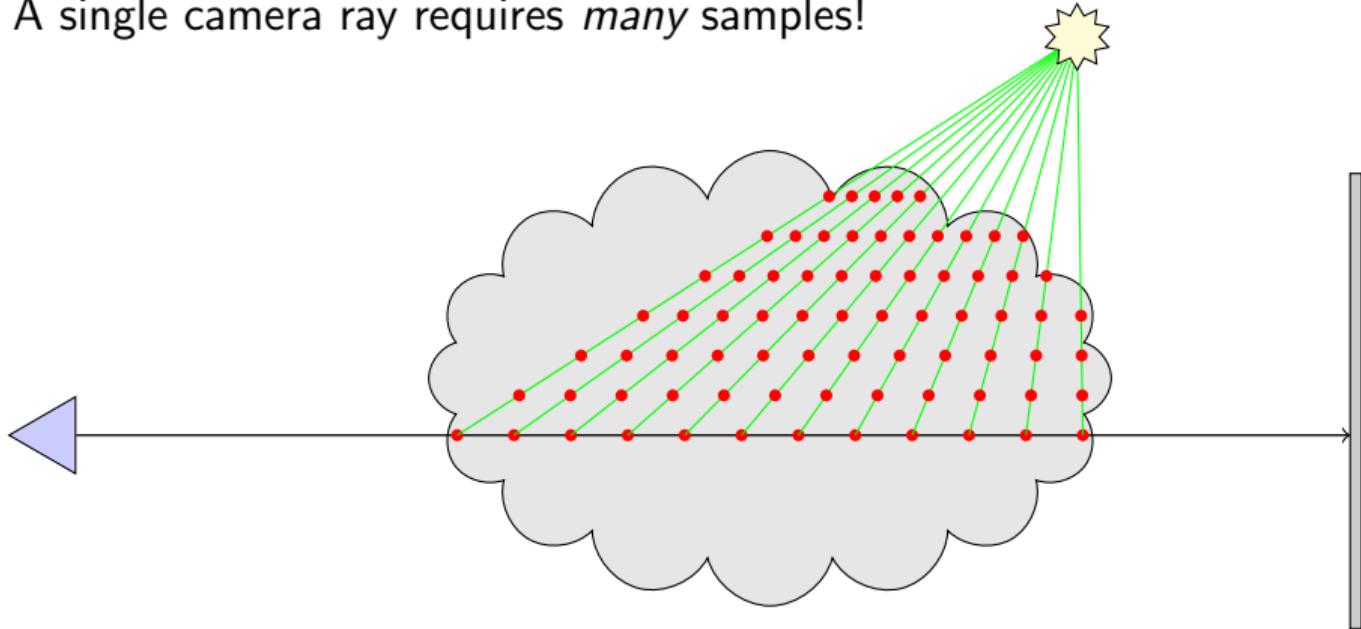
# Traditional Ray-marching

A single camera ray requires *many* samples!



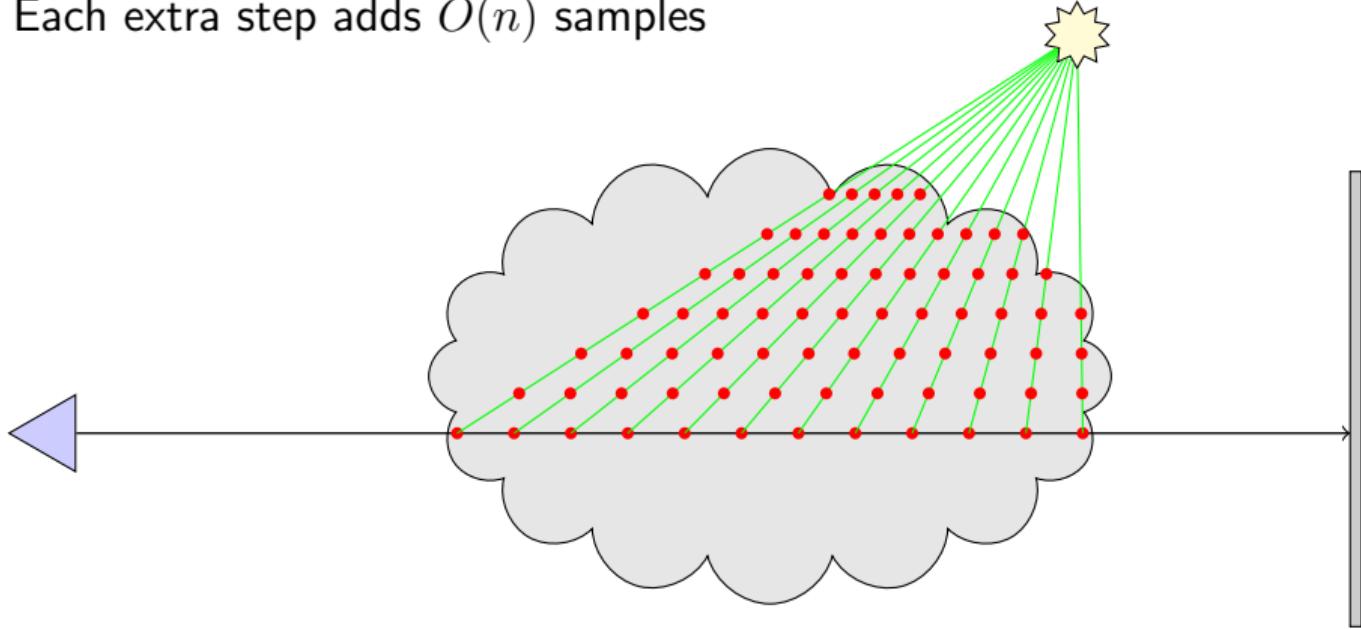
# Traditional Ray-marching

A single camera ray requires *many* samples!



# Traditional Ray-marching

Each extra step adds  $O(n)$  samples



## Previous Work: Biased Methods (1/2)

Computing lighting at every sample:

- ▶ Brute force algorithm is  $O(n^2)$

## Previous Work: Biased Methods (1/2)

Computing lighting at every sample:

- ▶ Brute force algorithm is  $O(n^2)$
- ▶ Requires use of cached lighting

## Previous Work: Biased Methods (1/2)

Computing lighting at every sample:

- ▶ Brute force algorithm is  $O(n^2)$
- ▶ Requires use of cached lighting
- ▶ “Deep Shadow Maps” [Lokovic and Veach, 2000]

## Previous Work: Biased Methods (1/2)

Computing lighting at every sample:

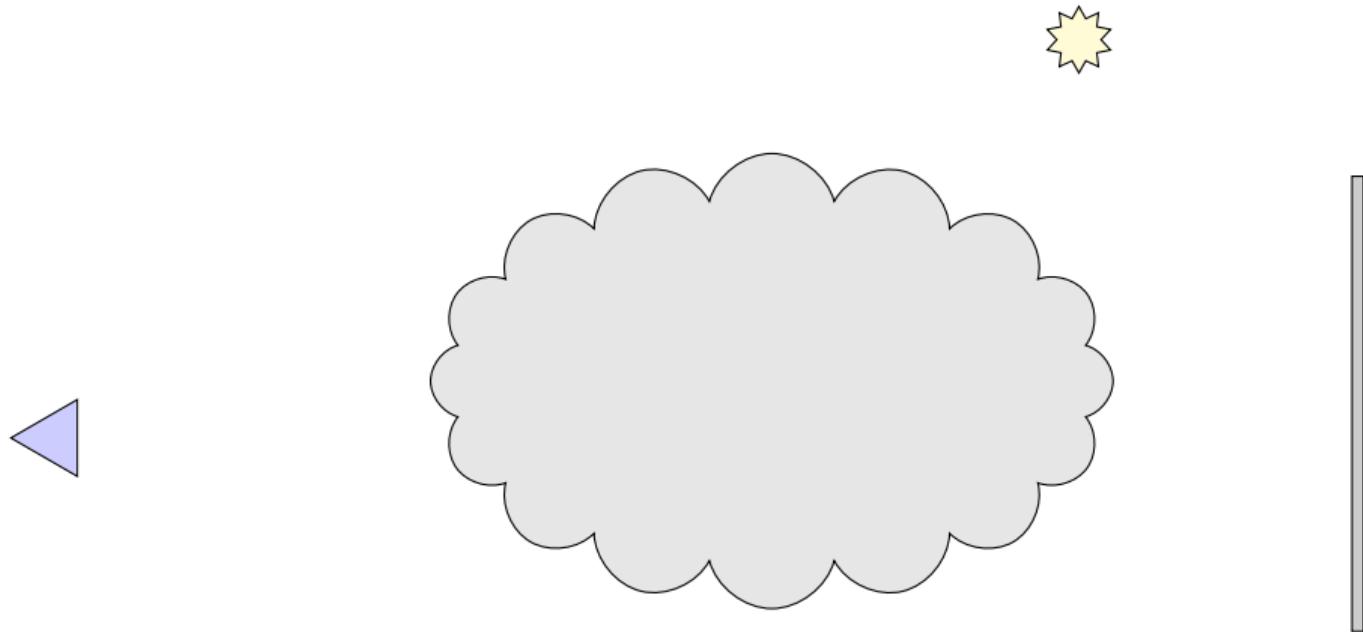
- ▶ Brute force algorithm is  $O(n^2)$
- ▶ Requires use of cached lighting
- ▶ “Deep Shadow Maps” [Lokovic and Veach, 2000]
- ▶ Multiple scattering requires further approximations

## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

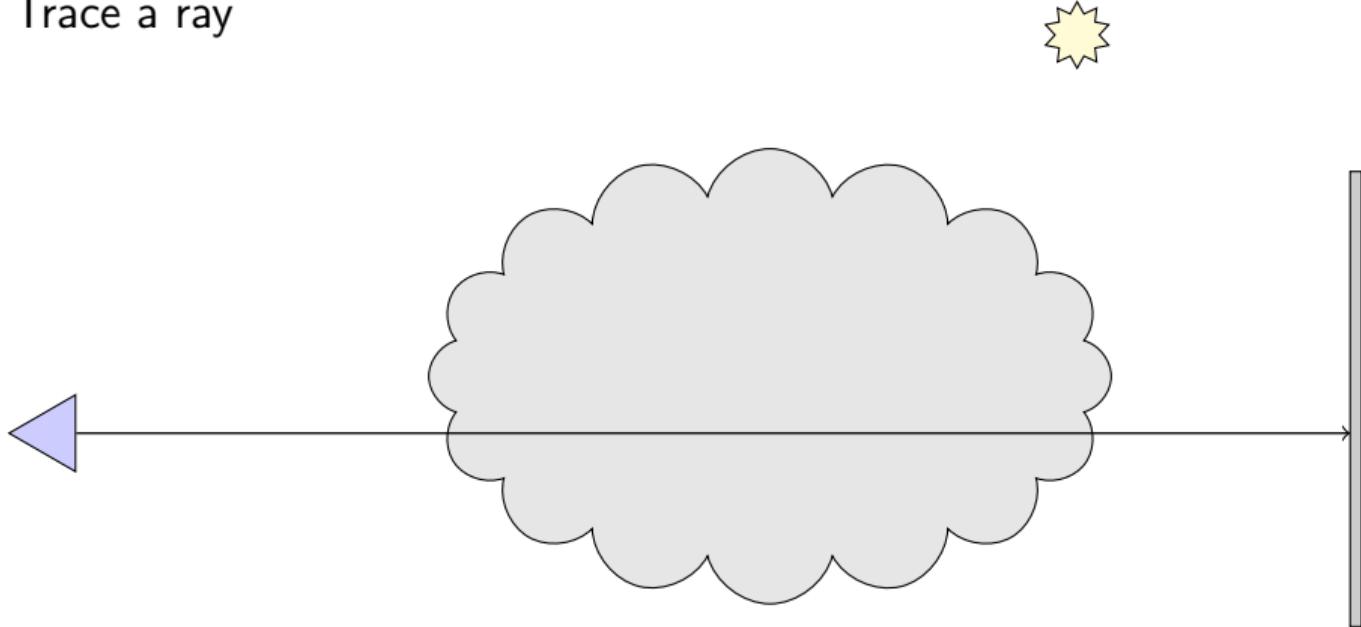
- ▶ “Rendering Participating Media with Bidirectional Path Tracing” [Lafortune et al, 1996]
- ▶ “Metropolis Light Transport for Participating Media” [Pauly et al, 2000]

# Stochastic Ray-marching



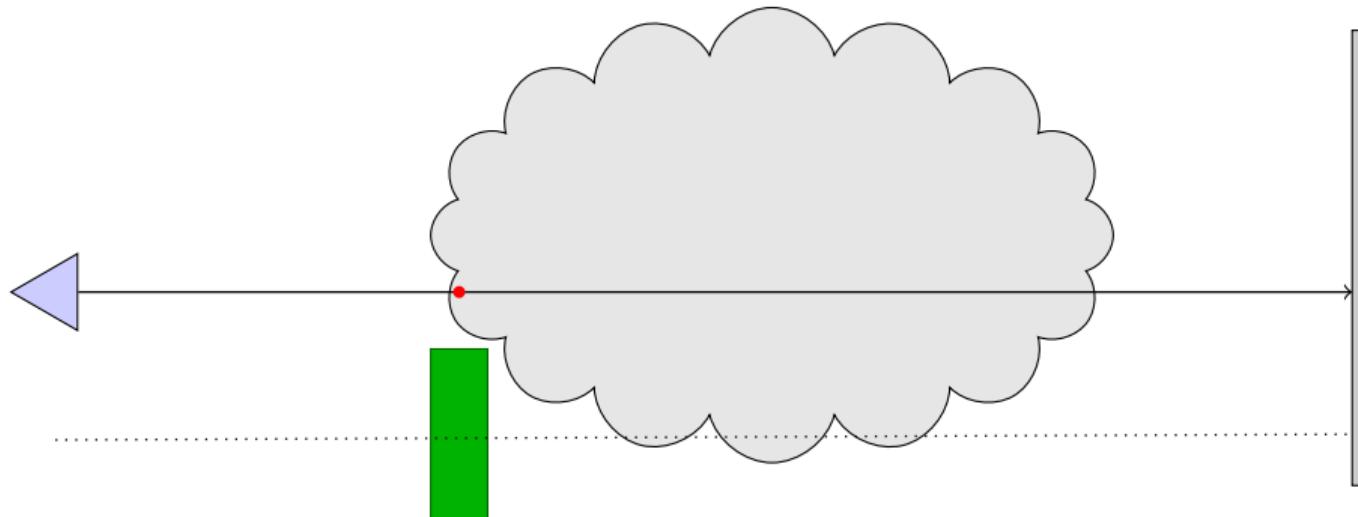
# Stochastic Ray-marching

Trace a ray



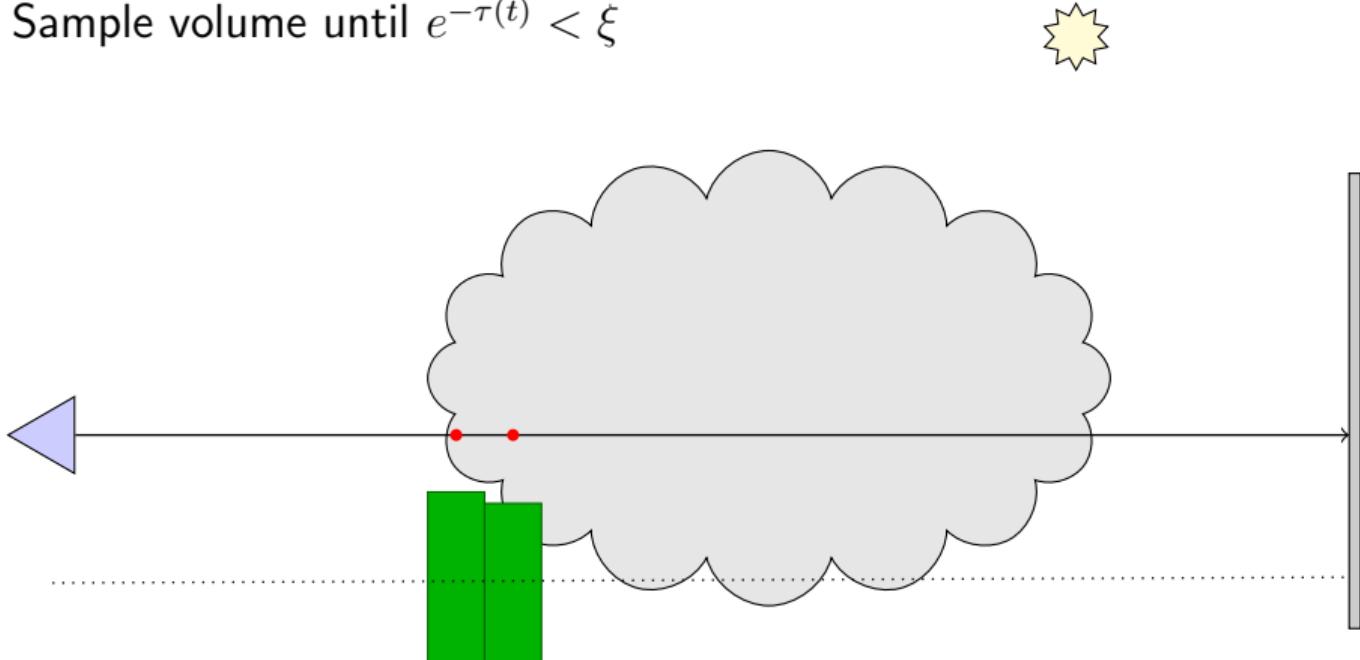
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



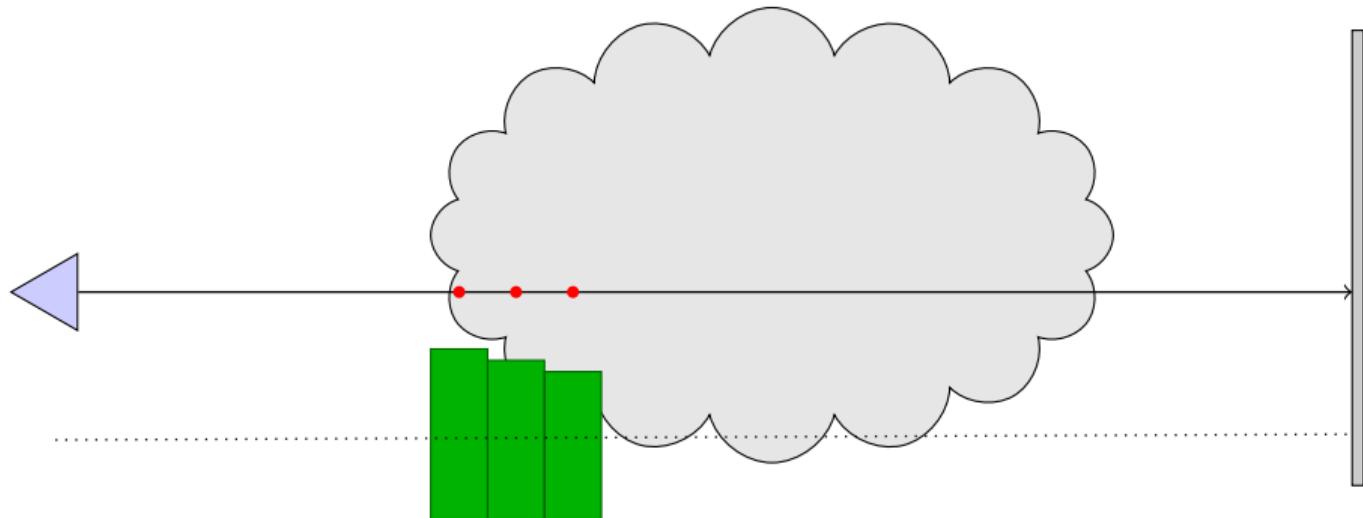
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



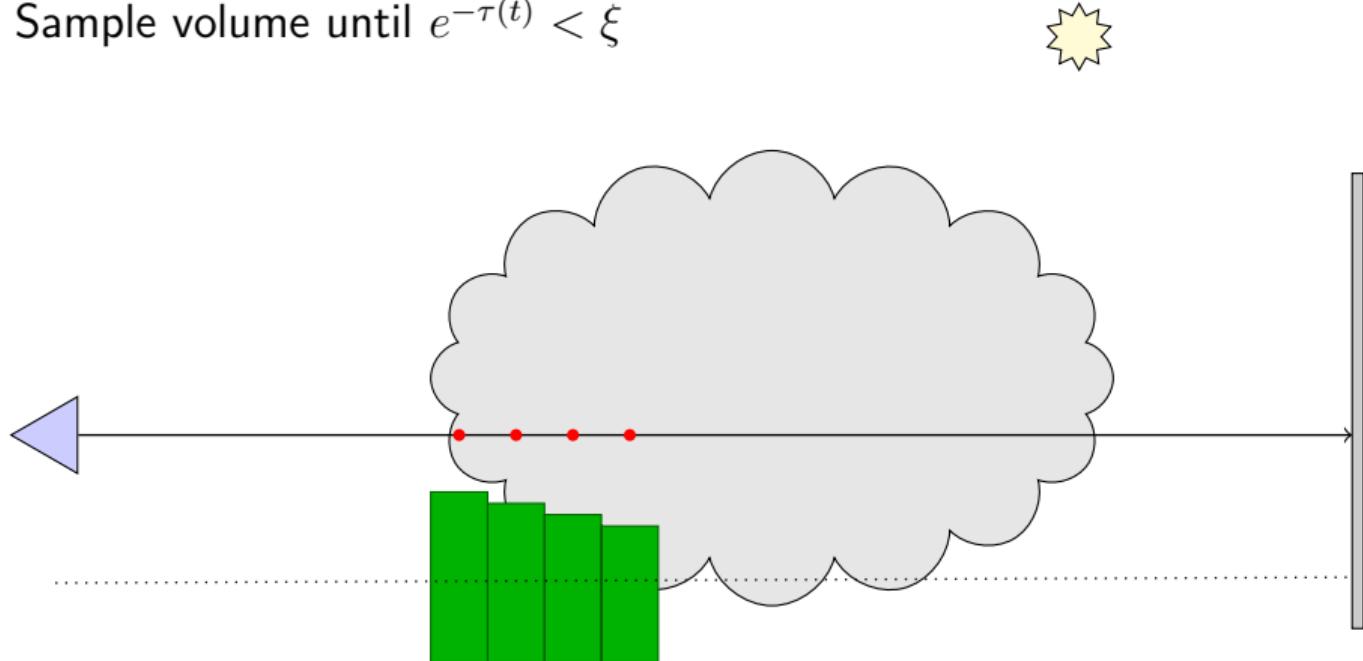
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



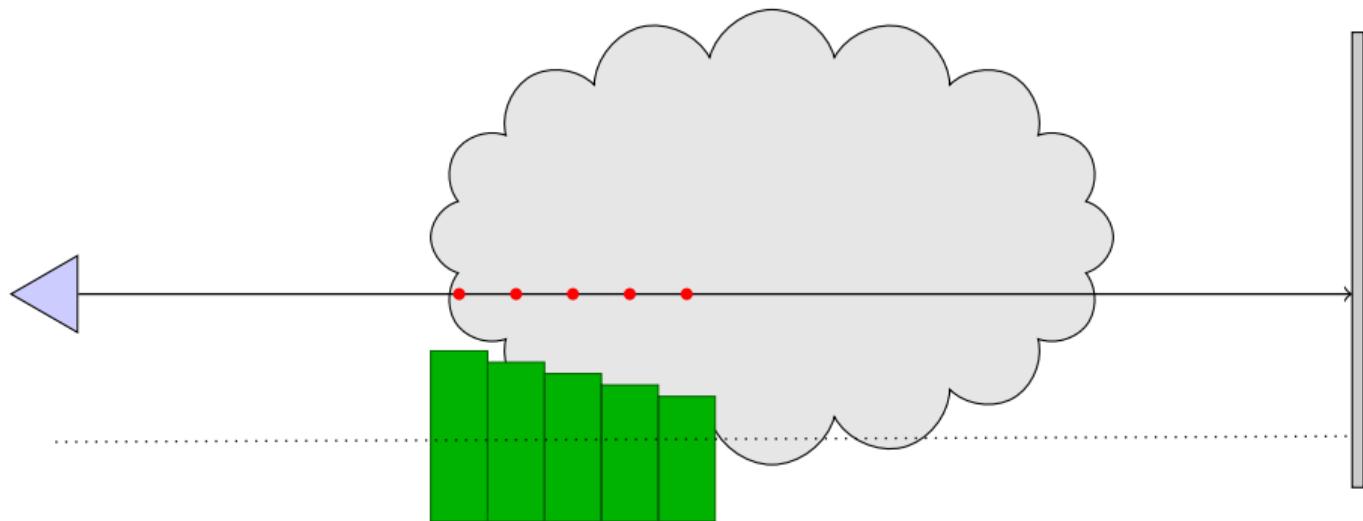
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



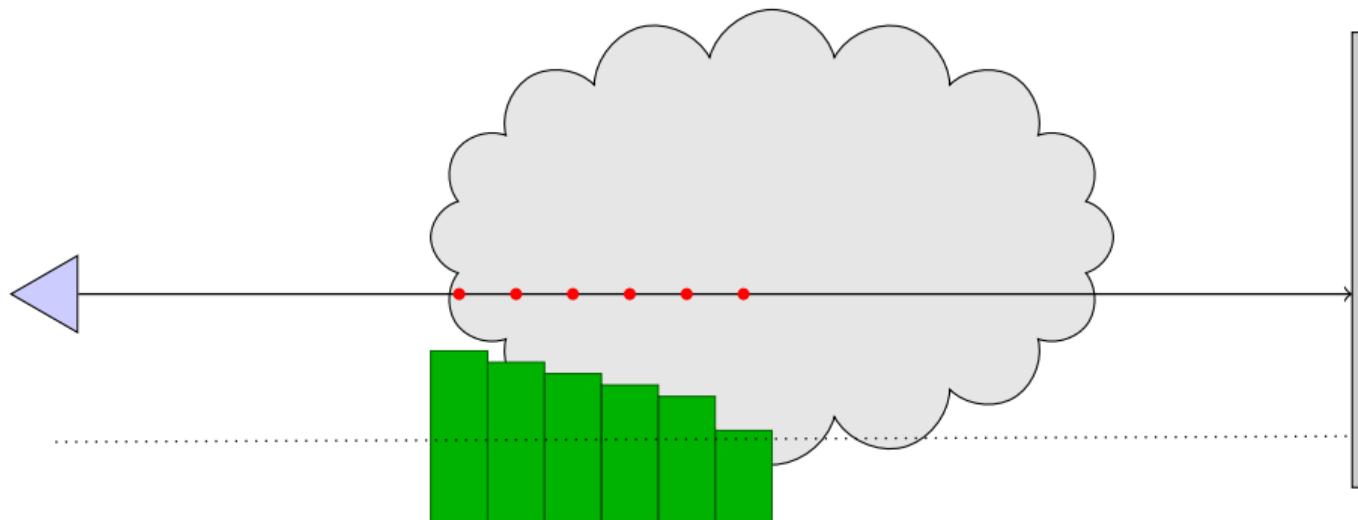
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



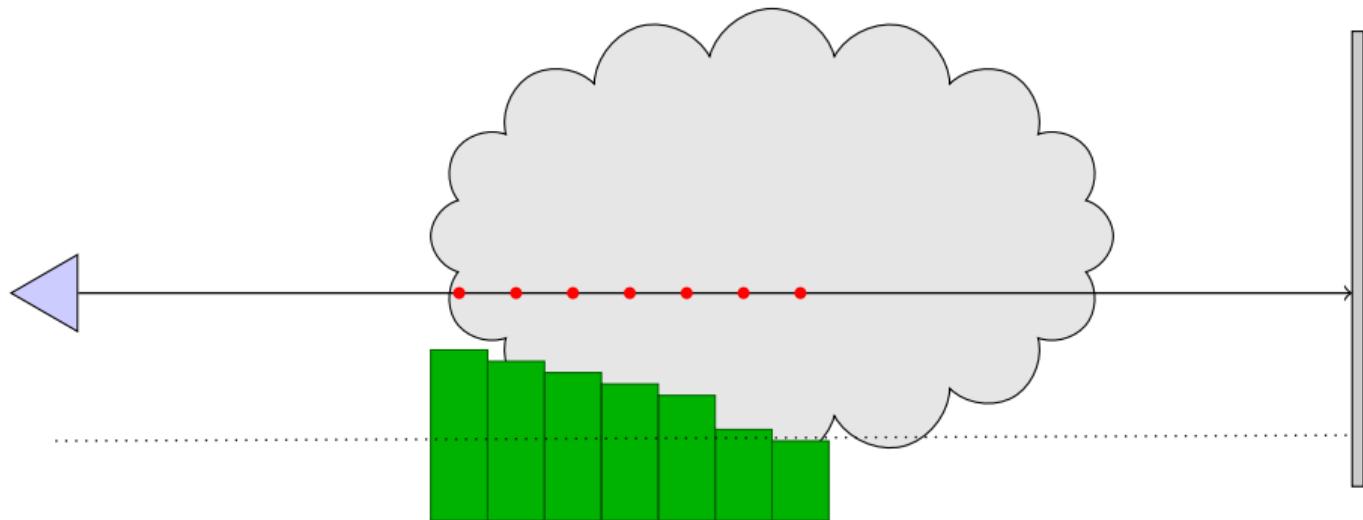
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



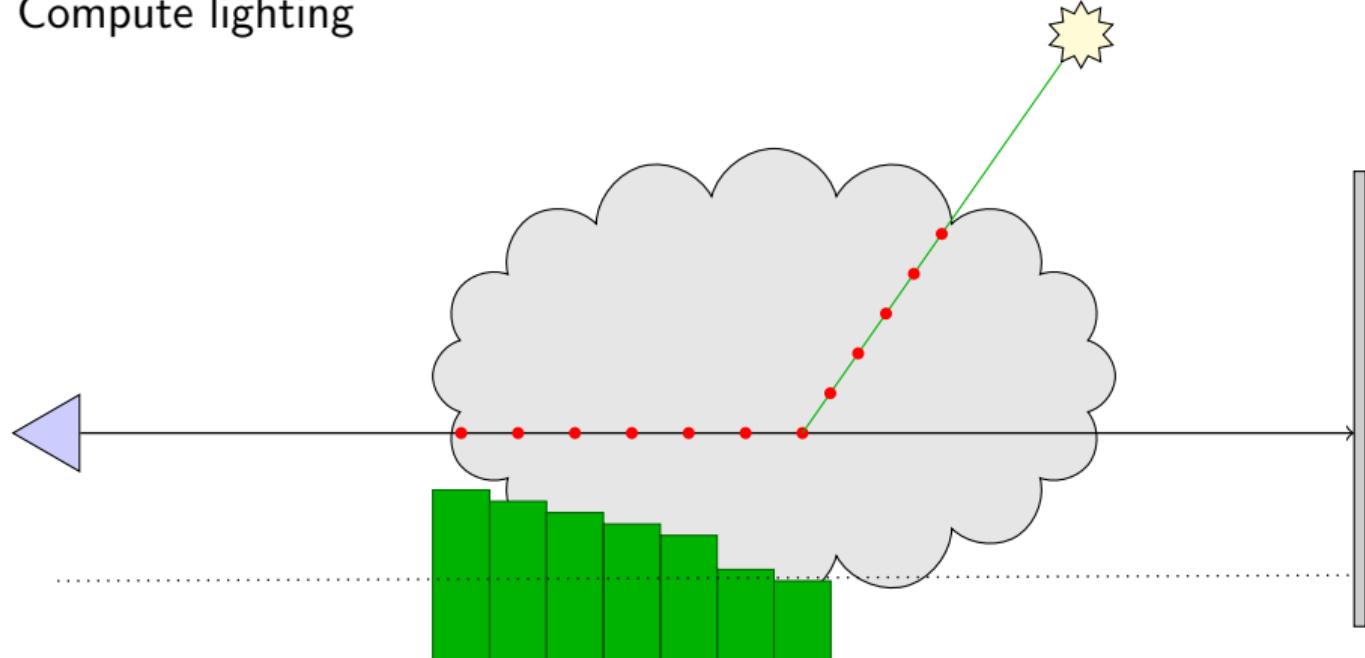
# Stochastic Ray-marching

Sample volume until  $e^{-\tau(t)} < \xi$



# Stochastic Ray-marching

Compute lighting



## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

- ▶ Similar to unbiased algorithm, but biased due to fixed size steps

## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

- ▶ Similar to unbiased algorithm, but biased due to fixed size steps
- ▶ Sample distribution is only implicit, pdf is not directly evaluable

## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

- ▶ Similar to unbiased algorithm, but biased due to fixed size steps
- ▶ Sample distribution is only implicit, pdf is not directly evaluable
- ▶ Taking  $n$  samples for a single ray requires restarting the ray march ...

## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

- ▶ Similar to unbiased algorithm, but biased due to fixed size steps
- ▶ Sample distribution is only implicit, pdf is not directly evaluable
- ▶ Taking  $n$  samples for a single ray requires restarting the ray march ...
- ▶ ... or sorting the sampling pattern ( $\xi_1 > \xi_2 > \dots > \xi_n$ )

## Previous Work: Biased Methods (2/2)

Using ray marching to propose a sample point:

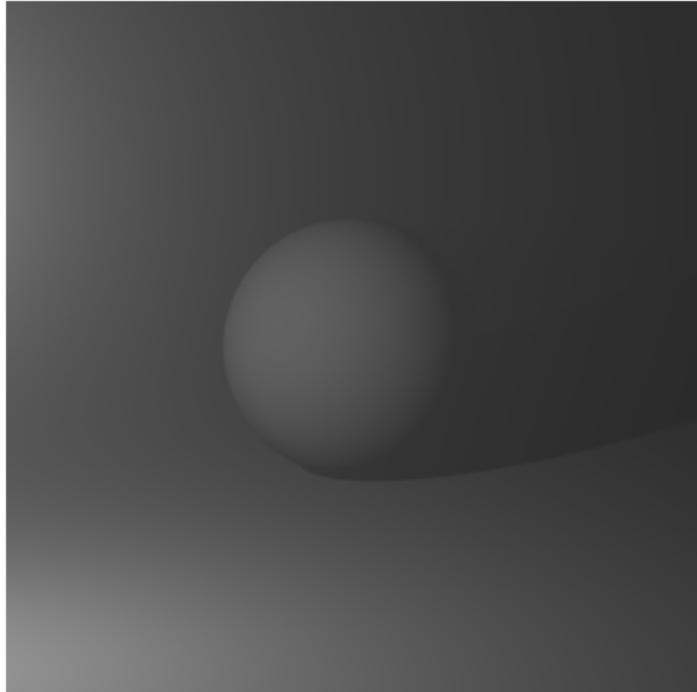
- ▶ Similar to unbiased algorithm, but biased due to fixed size steps
- ▶ Sample distribution is only implicit, pdf is not directly evaluable
- ▶ Taking  $n$  samples for a single ray requires restarting the ray march ...
- ▶ ... or sorting the sampling pattern ( $\xi_1 > \xi_2 > \dots > \xi_n$ )
- ▶ Lighting and transmission computed with the same number of samples

# Our Algorithm

- ▶ Decouples ray marching from lighting
- ▶ PDF along the ray can be evaluated and inverted efficiently
- ▶ Allows combining this work with methods from previous talk
- ▶ No restrictions on shaders
- ▶ No large data structures

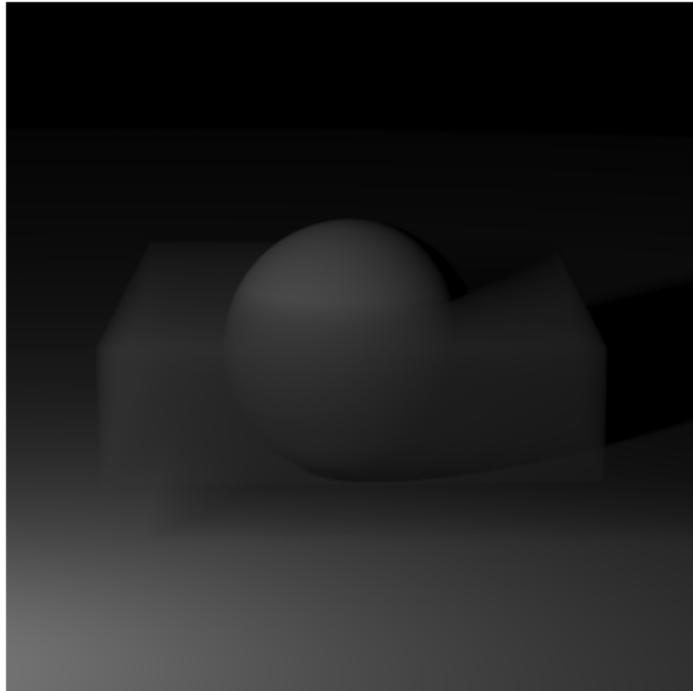
# Observation

- ▶ Homogeneous case is efficient



## Observation

- ▶ Homogeneous case is efficient
- ▶ Even when restricted to a finite region

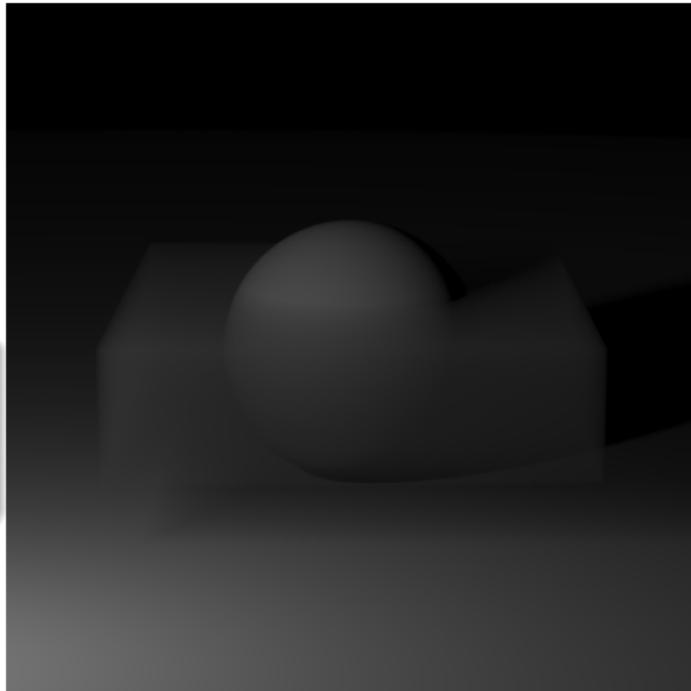


# Observation

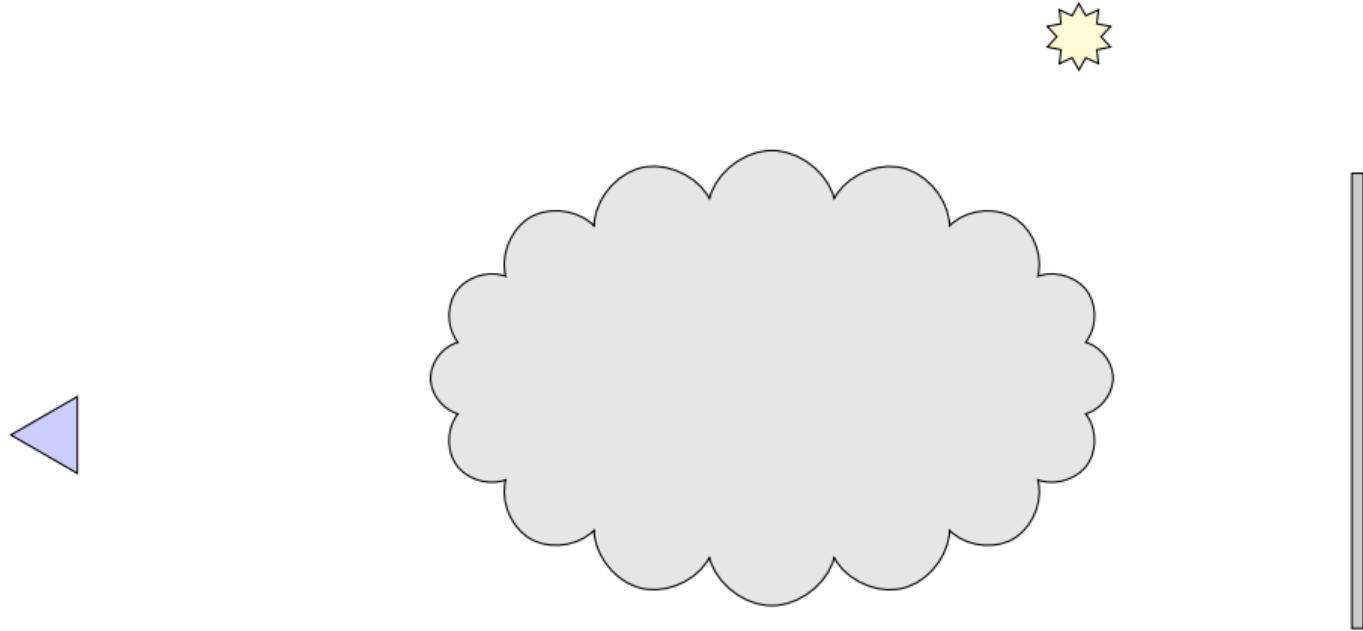
- ▶ Homogeneous case is efficient
- ▶ Even when restricted to a finite region

## Idea

Extend our methods to deal with  $n$  homogeneous segments along the ray

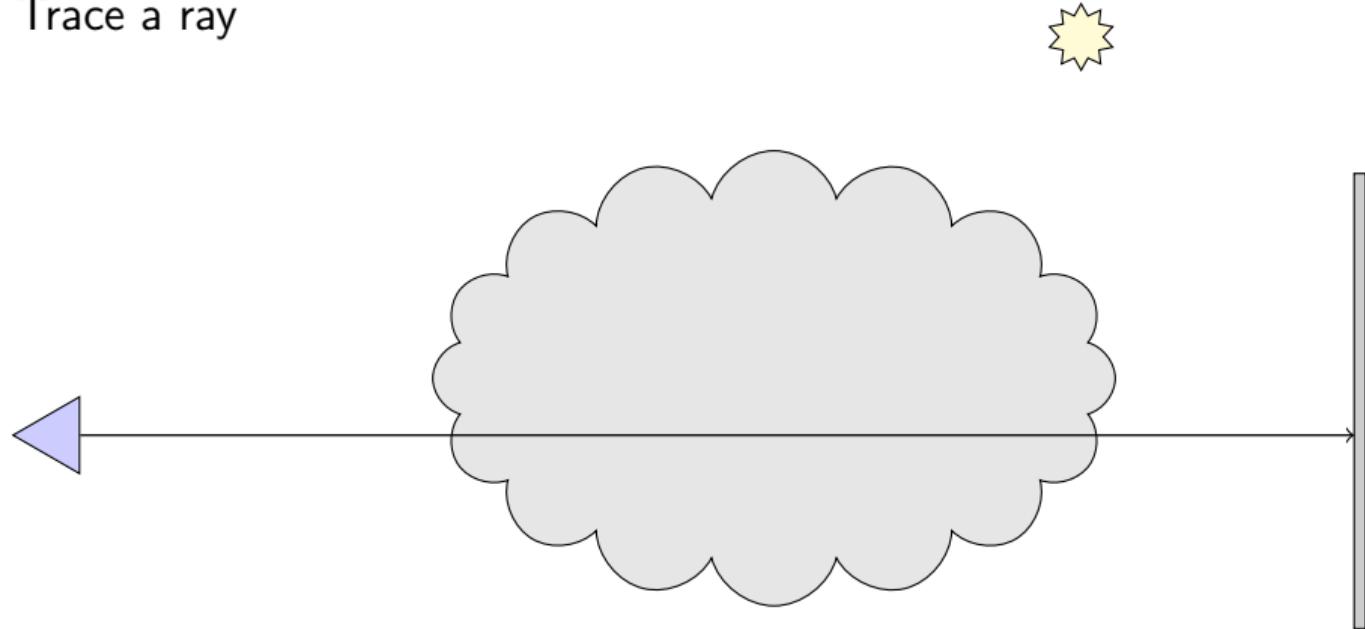


# Decoupled Ray-marching



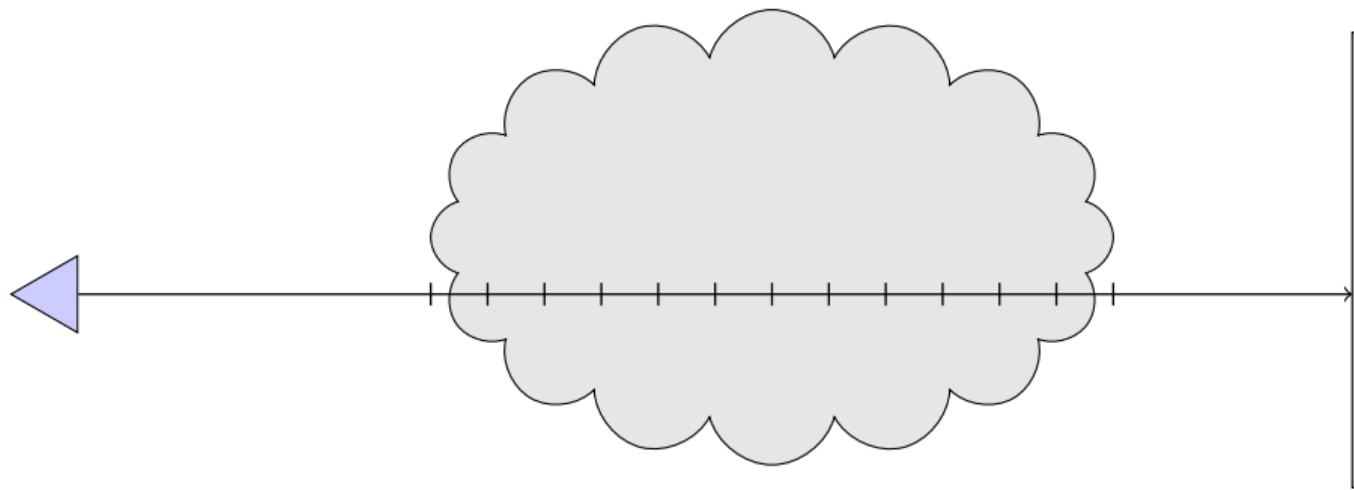
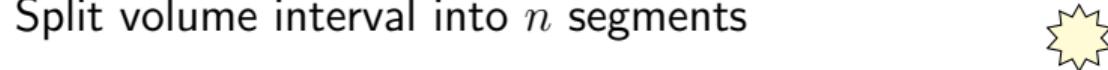
# Decoupled Ray-marching

Trace a ray



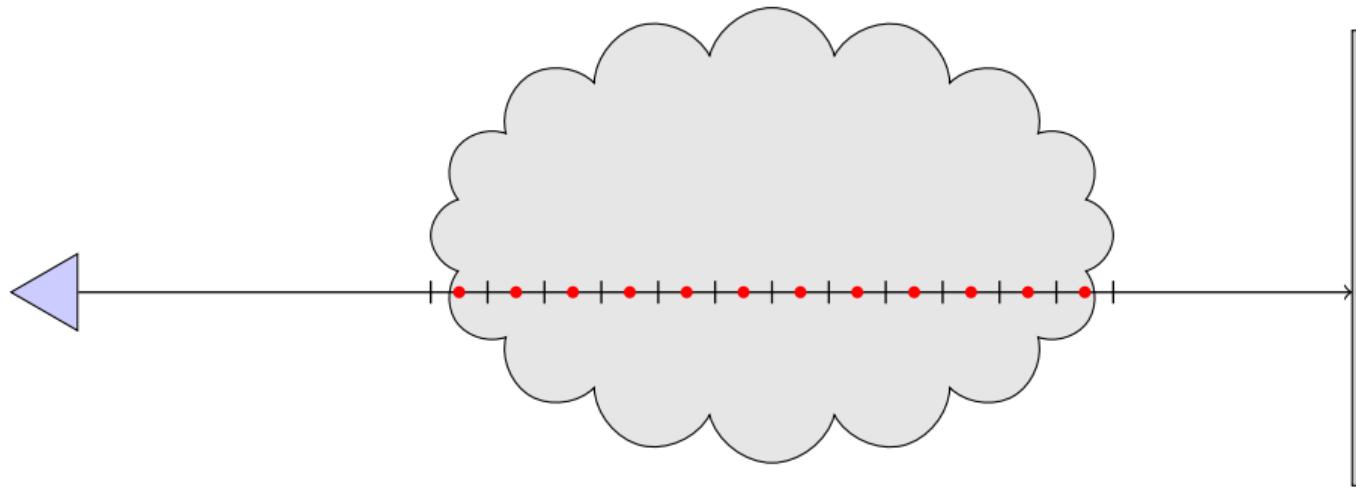
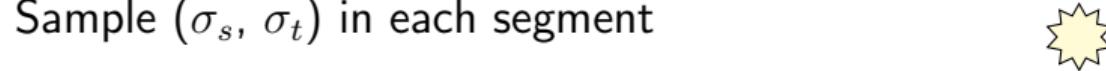
# Decoupled Ray-marching

Split volume interval into  $n$  segments



# Decoupled Ray-marching

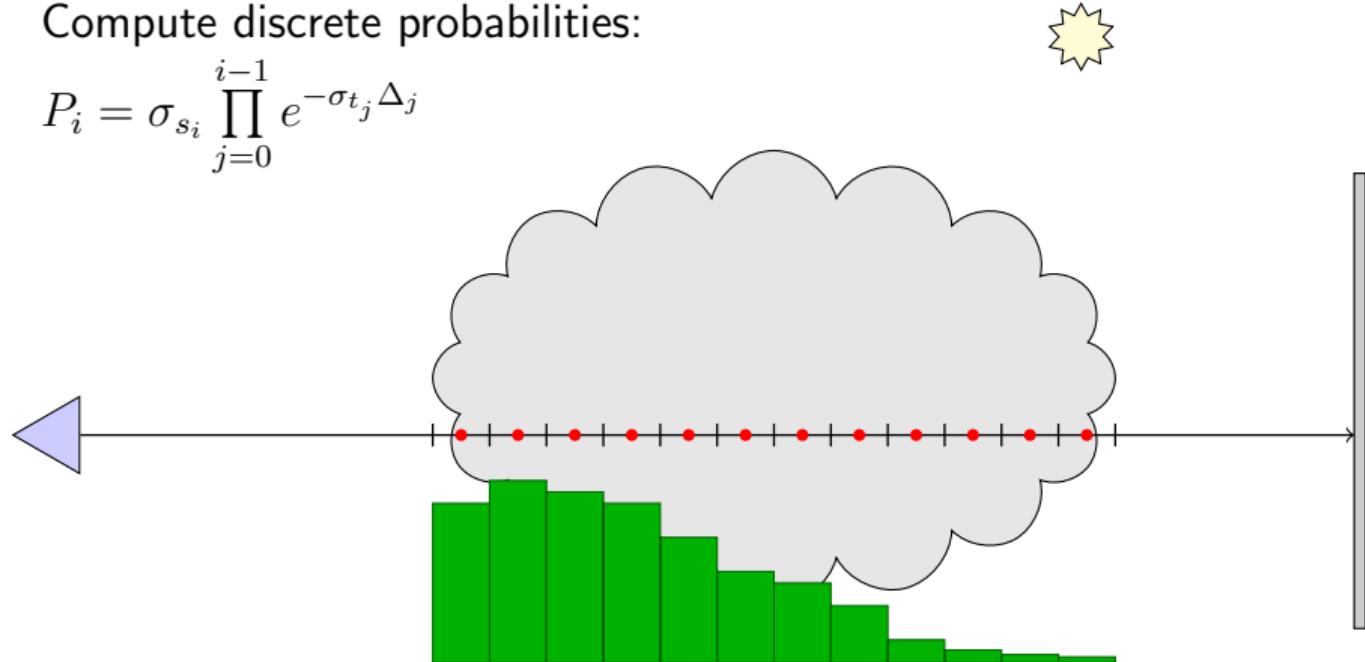
Sample  $(\sigma_s, \sigma_t)$  in each segment



# Decoupled Ray-marching

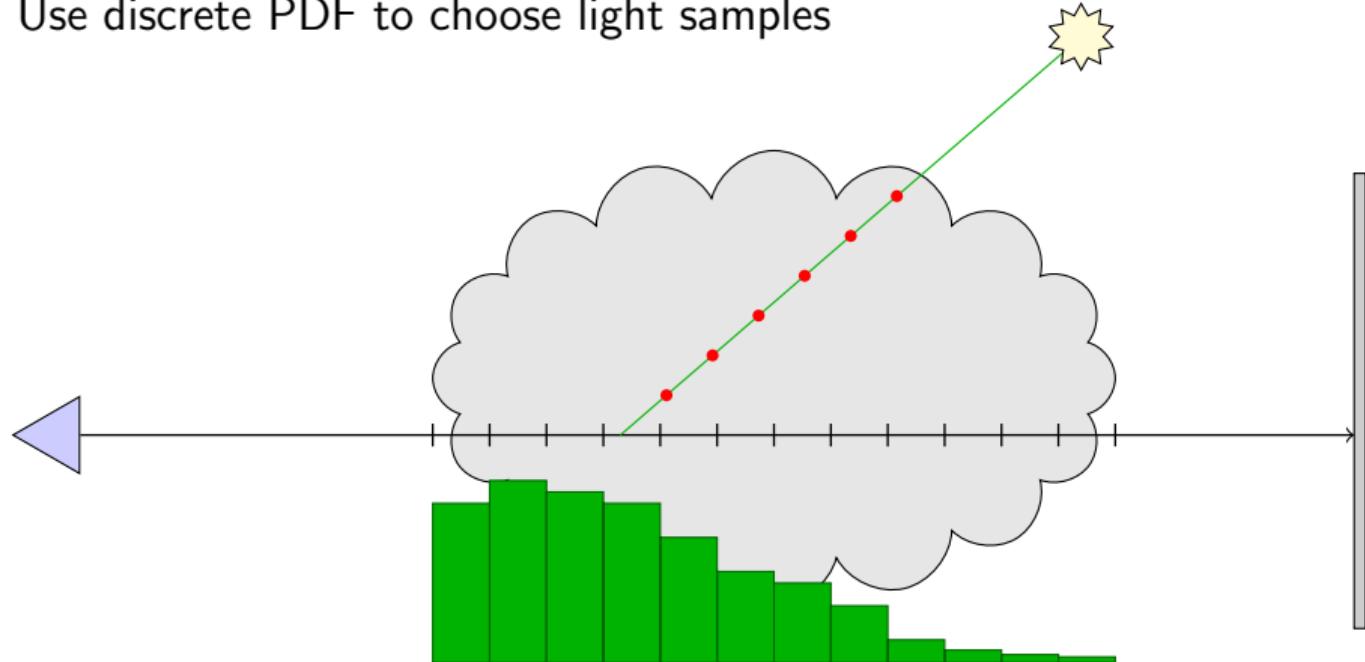
Compute discrete probabilities:

$$P_i = \sigma_{s_i} \prod_{j=0}^{i-1} e^{-\sigma_{t_j} \Delta_j}$$



# Decoupled Ray-marching

Use discrete PDF to choose light samples



## Decoupled Ray-marching: Advantages

- ▶ Unbiased when segments are really homogeneous

## Decoupled Ray-marching: Advantages

- ▶ Unbiased when segments are really homogeneous
- ▶ Converges as step size approaches Nyquist limit

## Decoupled Ray-marching: Advantages

- ▶ Unbiased when segments are really homogeneous
- ▶ Converges as step size approaches Nyquist limit
- ▶  $\tau(t)$  estimate is *decoupled* from number of lighting samples

## Decoupled Ray-marching: Advantages

- ▶ Unbiased when segments are really homogeneous
- ▶ Converges as step size approaches Nyquist limit
- ▶  $\tau(t)$  estimate is *decoupled* from number of lighting samples
- ▶ Lighting samples focus on areas where  $\sigma_s > 0$

# Time Complexity

Assuming we take  $O(n)$  steps for ray-marching:

- ▶ Obtain CDF from PDF in  $O(n)$

# Time Complexity

Assuming we take  $O(n)$  steps for ray-marching:

- ▶ Obtain CDF from PDF in  $O(n)$
- ▶ Can draw samples in  $O(\log(n))$  (binary search on CDF)

# Time Complexity

Assuming we take  $O(n)$  steps for ray-marching:

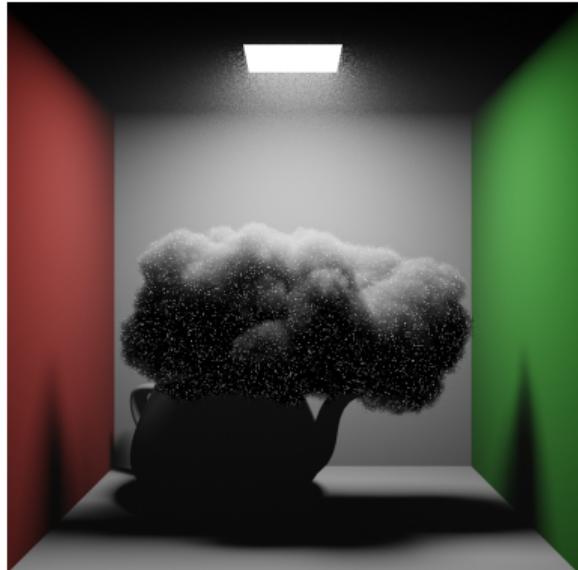
- ▶ Obtain CDF from PDF in  $O(n)$
- ▶ Can draw samples in  $O(\log(n))$  (binary search on CDF)
- ▶ Can evaluate PDF and  $\tau$  for any  $t$  in  $O(\log(n))$  (binary search on segments)

# Time Complexity

Assuming we take  $O(n)$  steps for ray-marching:

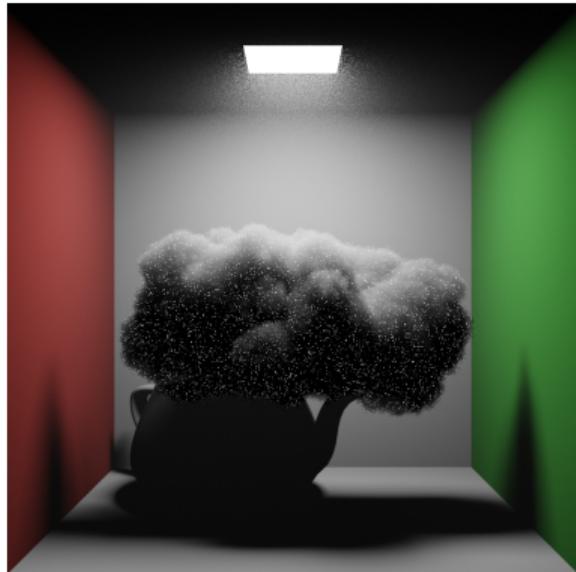
- ▶ Obtain CDF from PDF in  $O(n)$
- ▶ Can draw samples in  $O(\log(n))$  (binary search on CDF)
- ▶ Can evaluate PDF and  $\tau$  for any  $t$  in  $O(\log(n))$  (binary search on segments)
- ▶ Allows combinations of line-sampling strategies!

# Multiple Importance Sampling

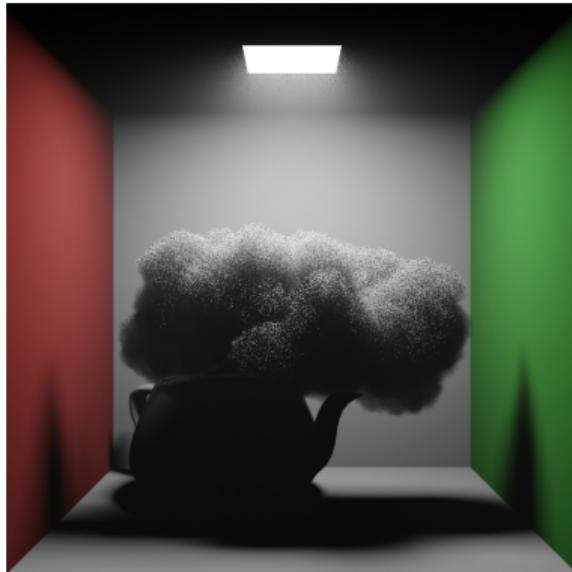


Density Sampling

# Multiple Importance Sampling

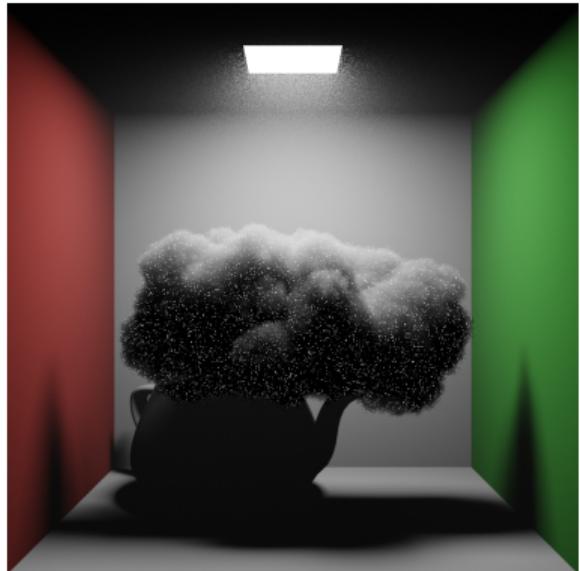


Density Sampling

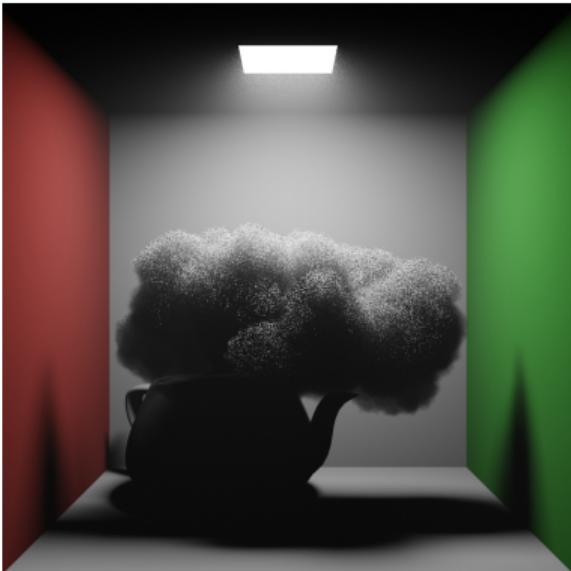


Equi-angular Sampling

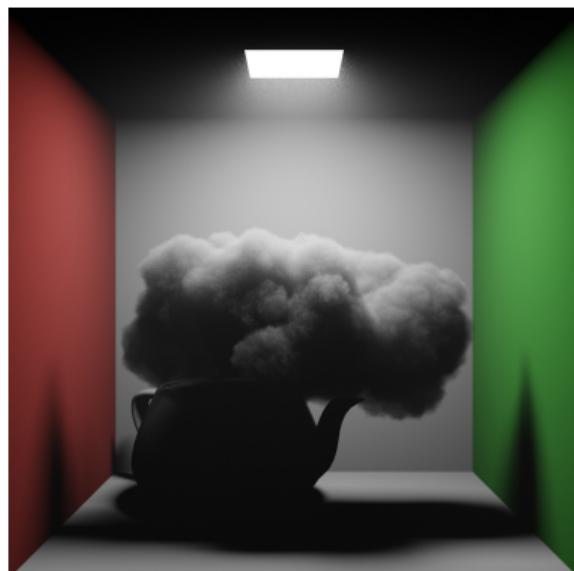
# Multiple Importance Sampling



Density Sampling

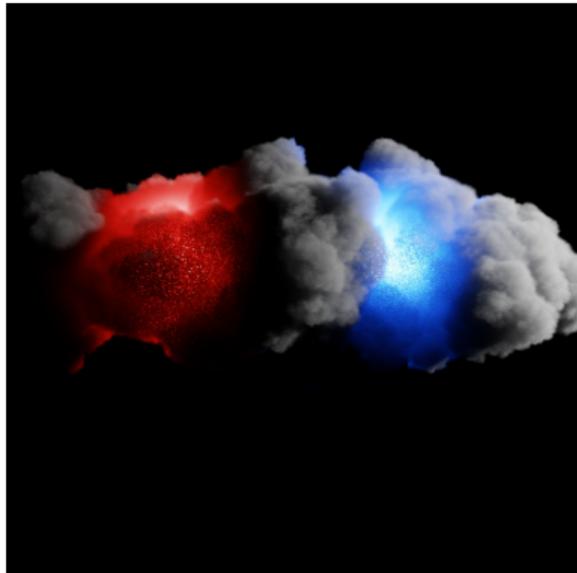


Equi-angular Sampling



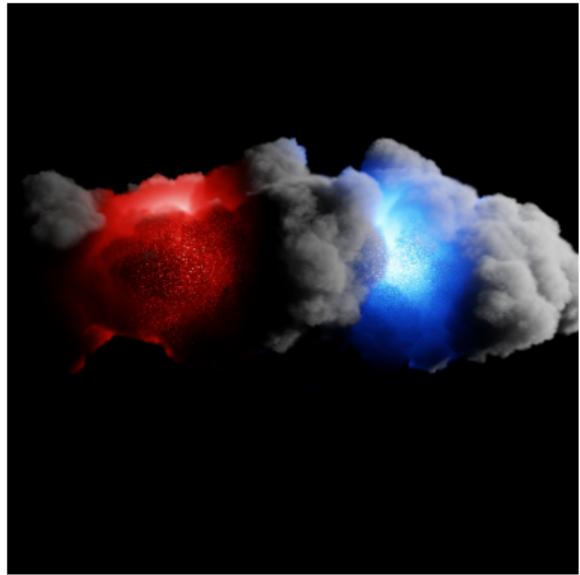
MIS

# Multiple Importance Sampling

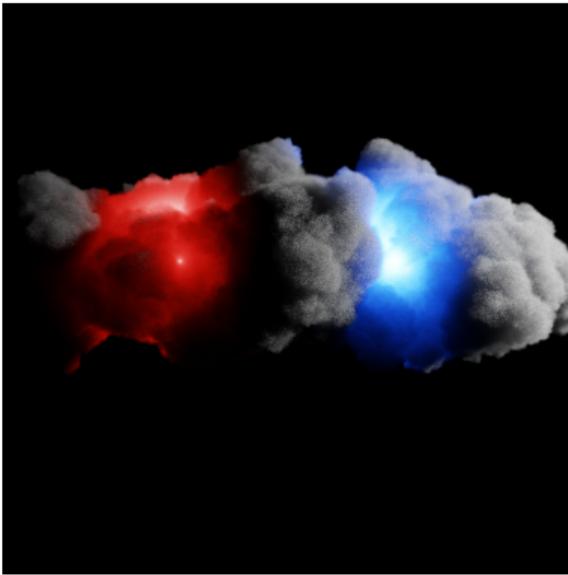


Density Sampling

# Multiple Importance Sampling

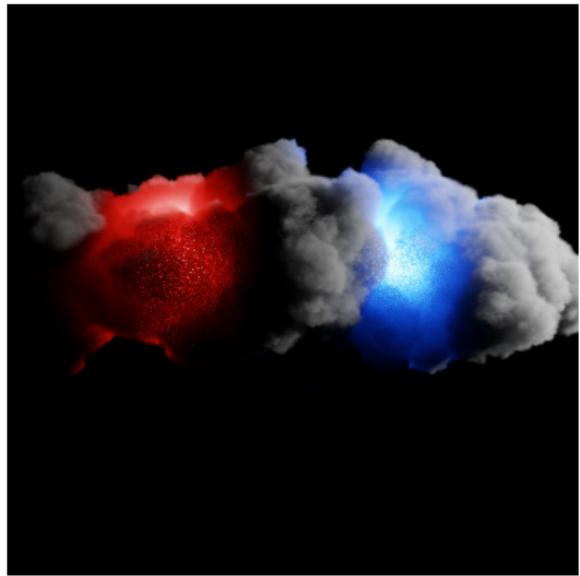


Density Sampling

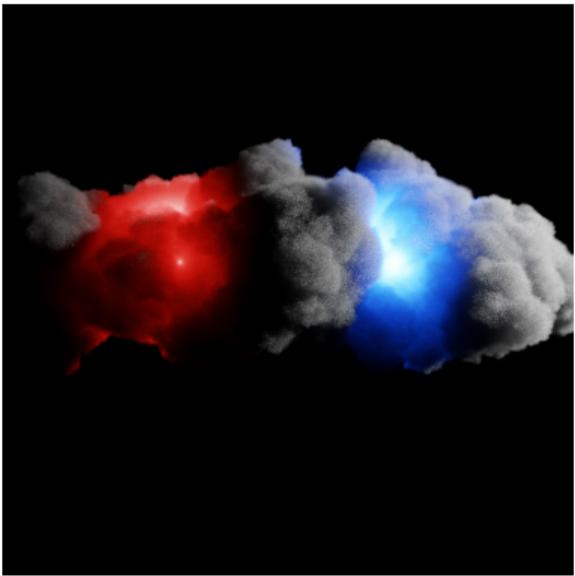


Equi-angular Sampling

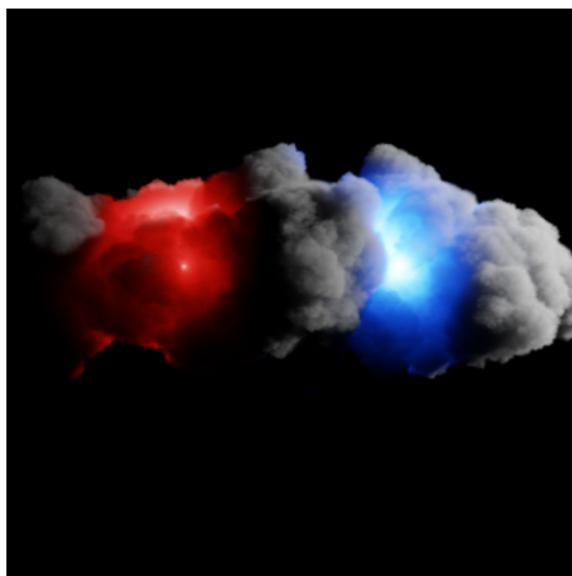
# Multiple Importance Sampling



Density Sampling



Equi-angular Sampling

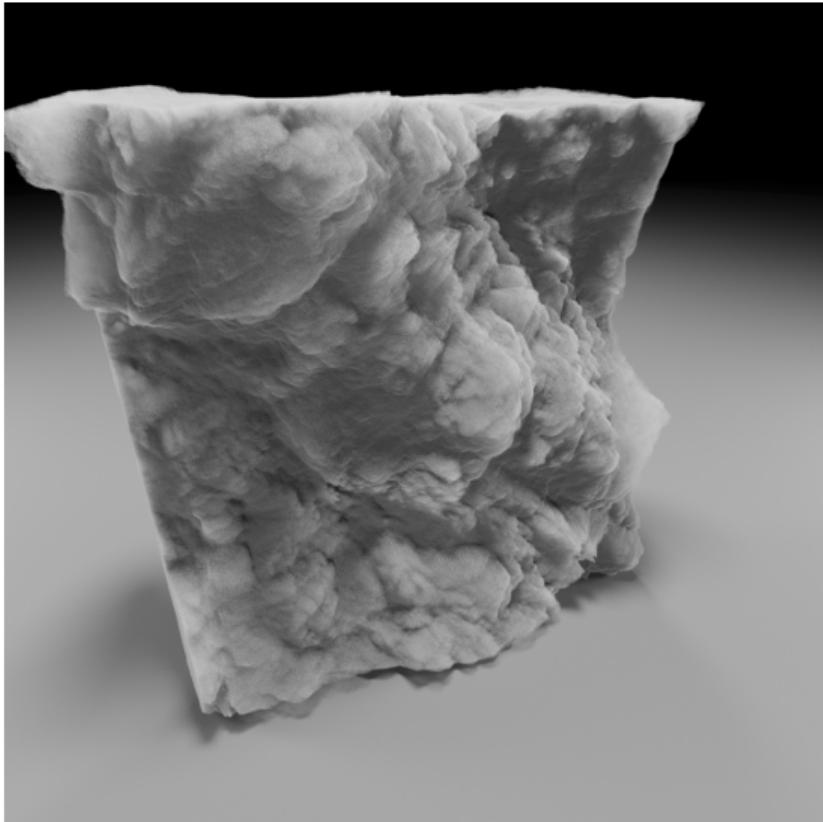
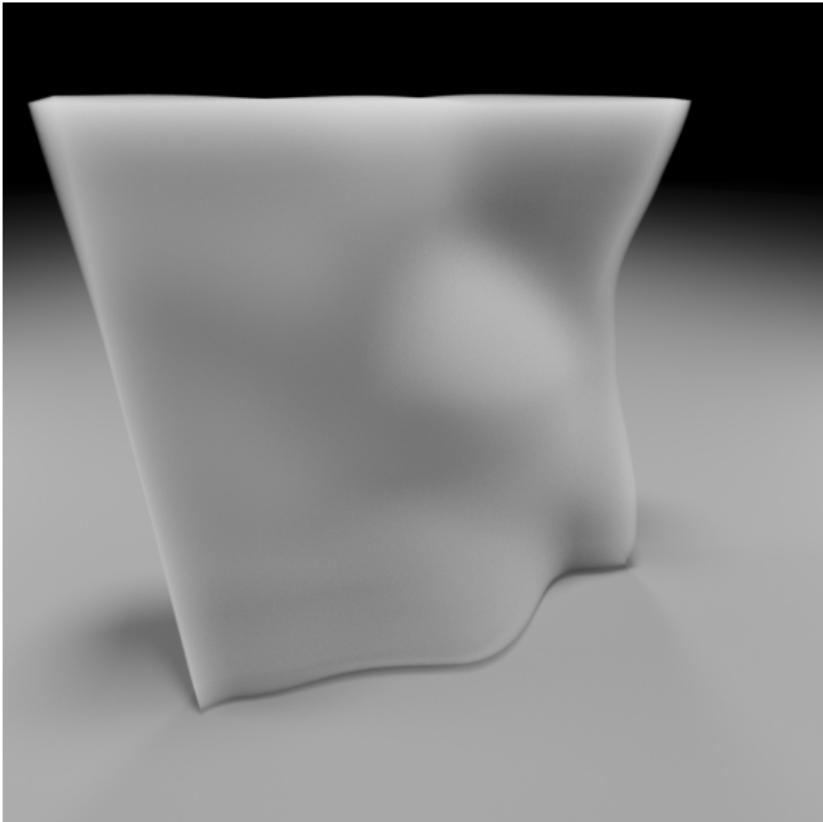


MIS

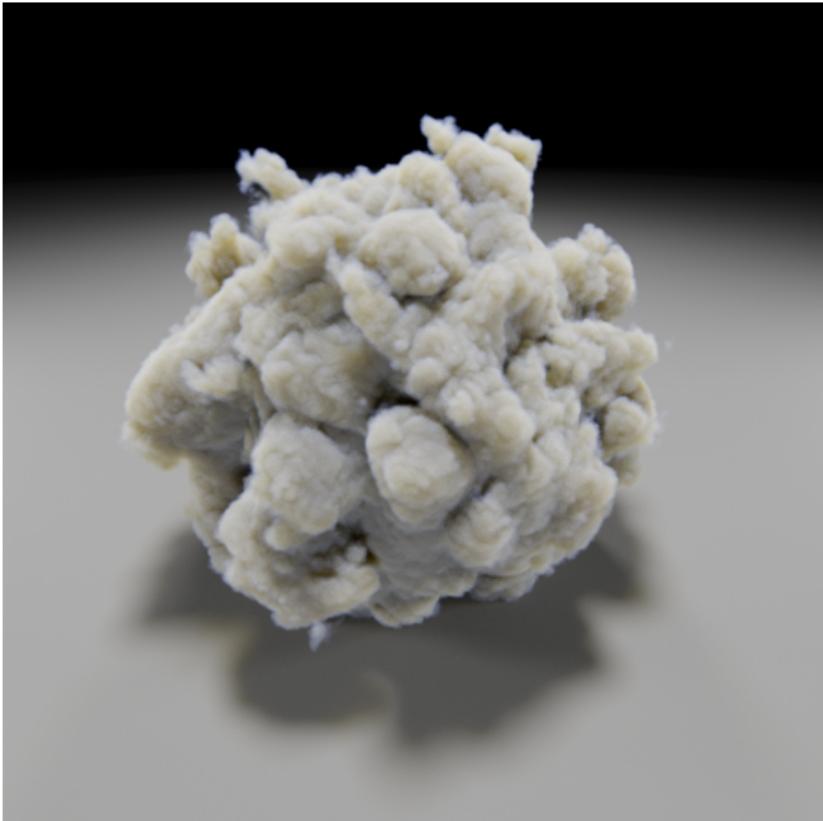
## Implementation Details: Memory Management

- ▶ Algorithm requires temporary arrays
- ▶ Use a per-thread memory pool with a stack interface
- ▶ Push pool state before a ray is processed
- ▶ Pop state when the ray volume shading is done
- ▶ Allows memory to grow only with the recursion depth (only a few Kb in practice)

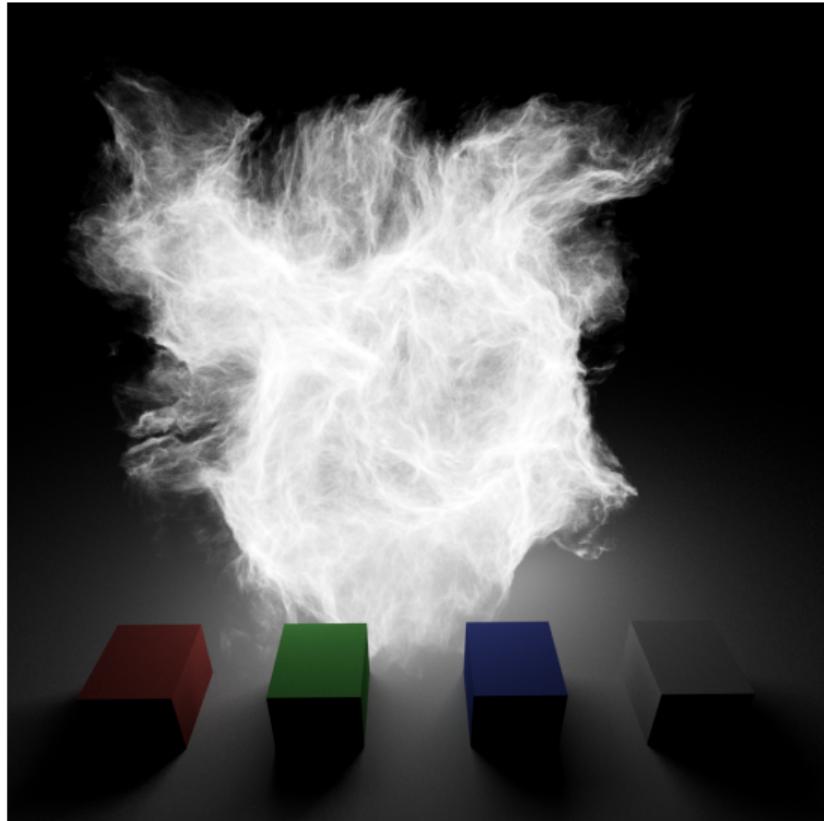
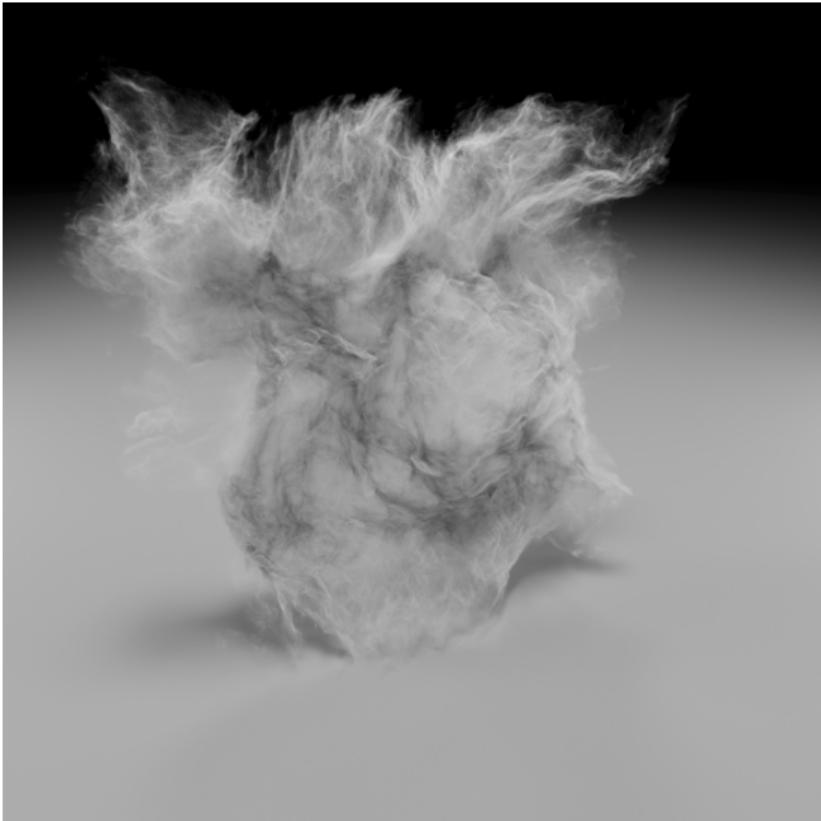
# Examples



## Examples



# Examples



# Examples



0 bounces  
3m28s



1 bounces  
10m14s



2 bounces  
20m21s

# Examples

Videos

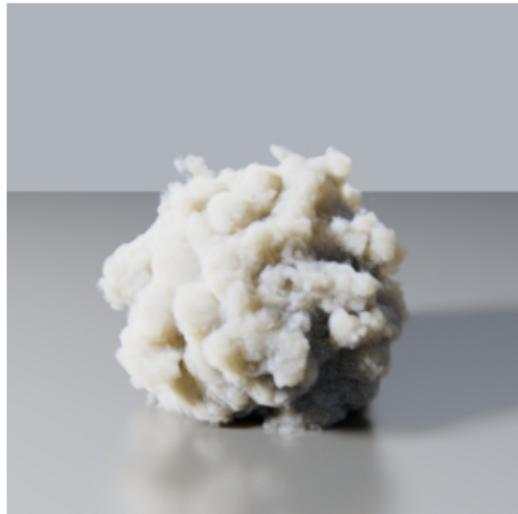
## Summary

- ▶ Practical method for ray traced lighting of heterogeneous volumes
- ▶ Lighting calculations are decoupled from volume sampling
- ▶ Supports procedural shaders
- ▶ Supports all light source types
- ▶ Allows for multiple scattering
- ▶ Does not require any (permanent) auxiliary data structures

## Future Work

- ▶ Investigate better weights for MIS
- ▶ Investigate sampling distributions for multiple scattering
- ▶ Automate the step size using ray differentials
- ▶ Efficient illumination from volumes (fire as light source)
- ▶ Out-of-core datasets

Thanks for listening!



Acknowledgment: Magnus Wrenninge (models, fluid sims)